

# **CSE 120**

# **Principles of Operating Systems**

**Fall 2021**

**Lecture 9: Memory Management Overview**

Yiying Zhang

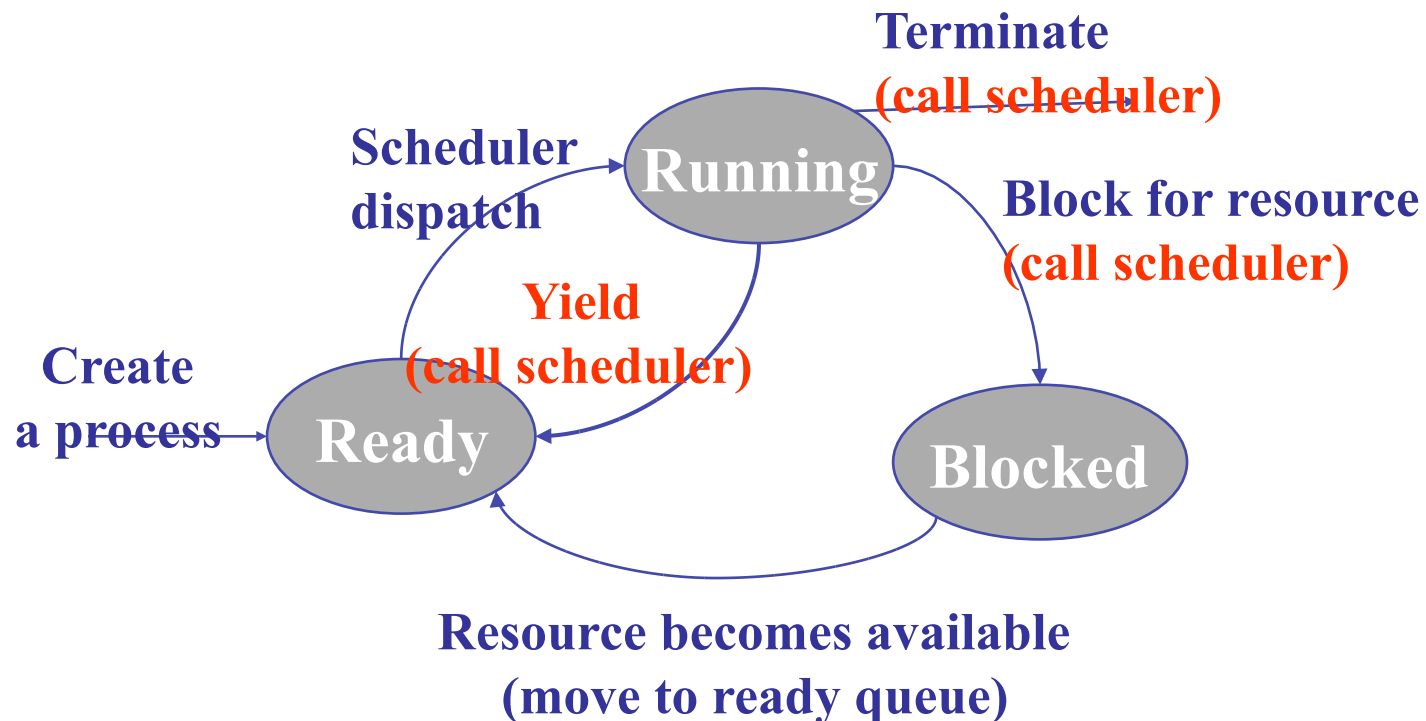
# Announcements

---

- Homework 2 due tonight
- Start working on Project 2!
  - ♦ It will be harder than project 1 and need more time

# [lec8] Non-Preemptive Scheduling

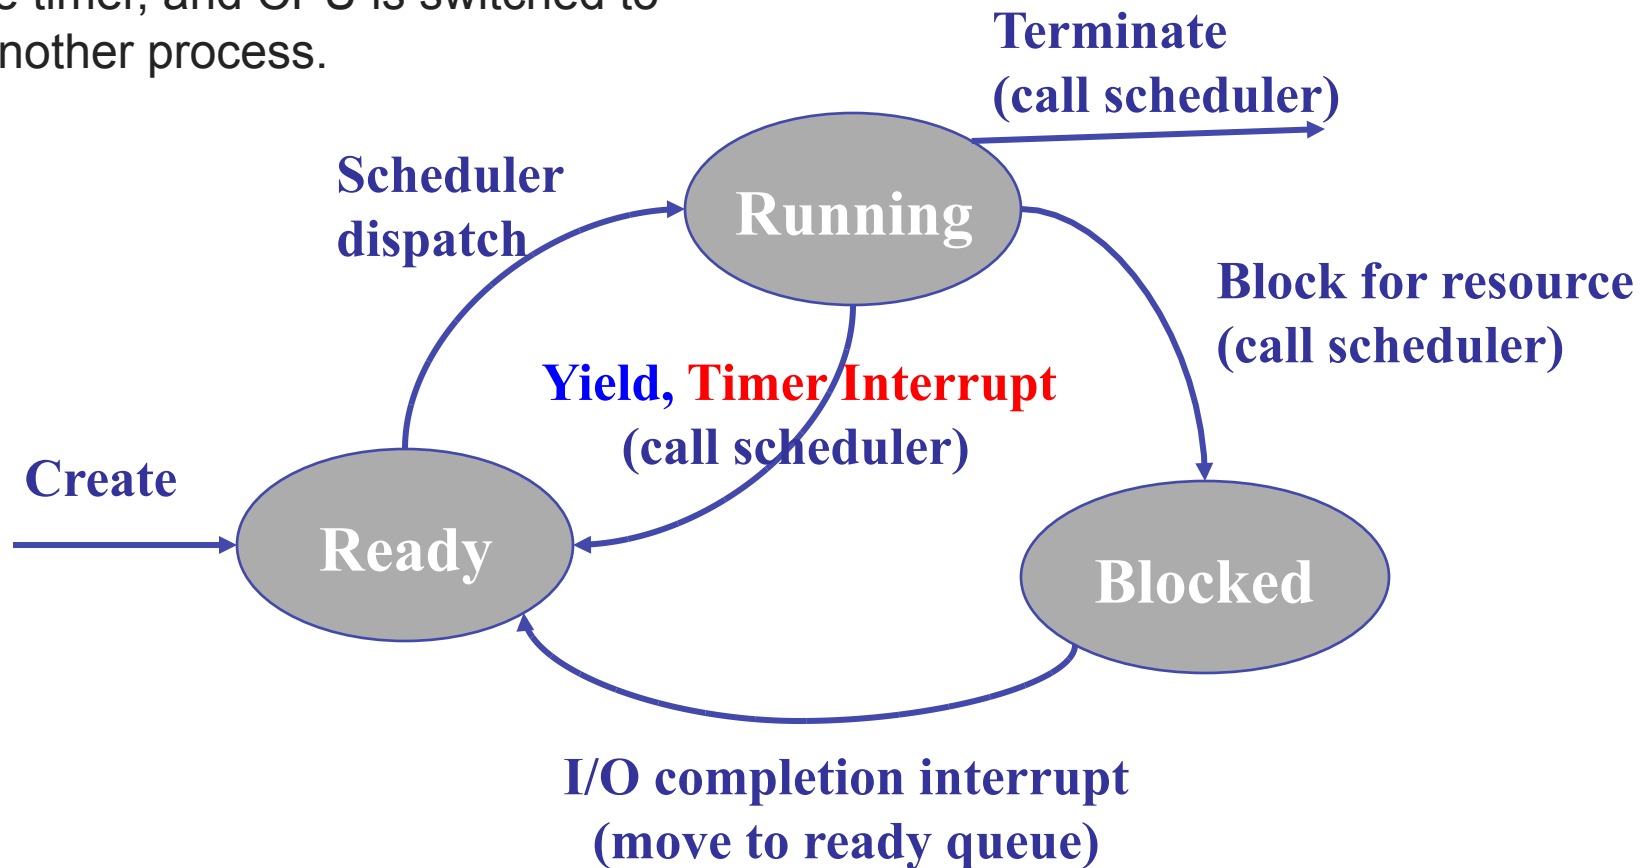
- OS only has a chance to schedule threads on a core when the current running thread leaves its running state
  - ♦ Yield, terminate, blocked by I/O, etc.



- How can we force a thread off its running state?

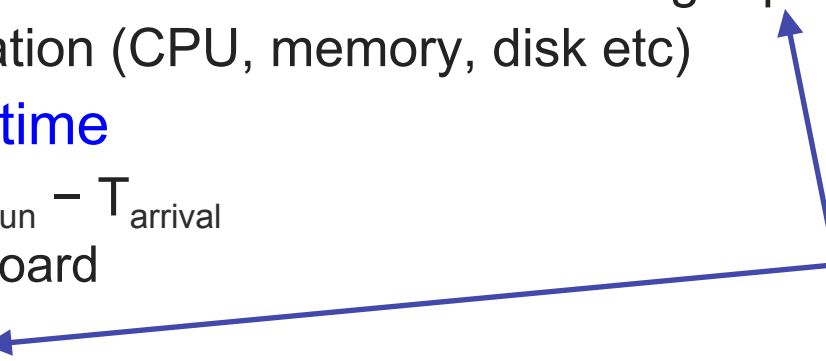
# [lec8] Preemptive Scheduling

A running process is interrupted by the timer, and CPU is switched to run another process.



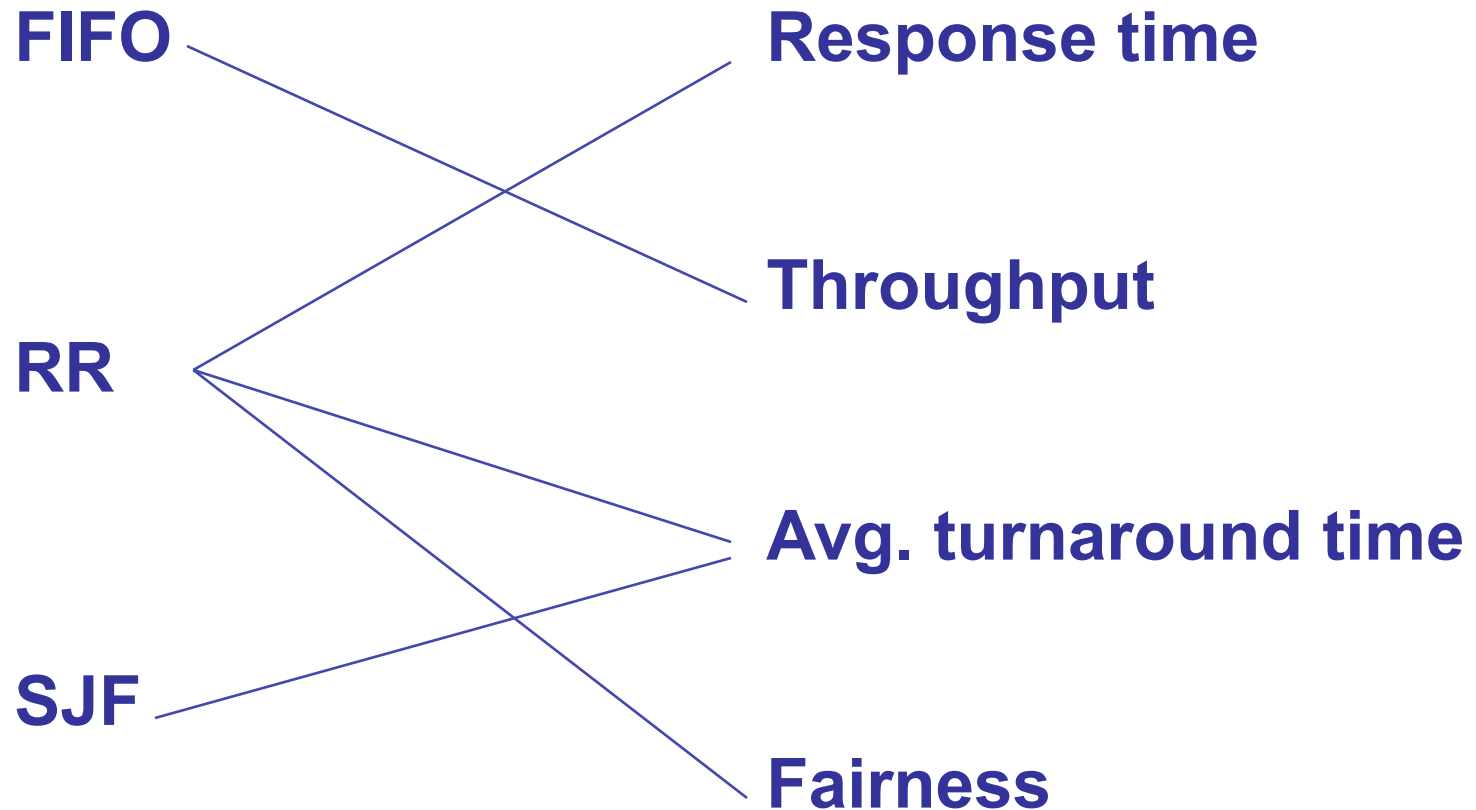
# [lec8] Goals and Assumptions

---

- Goals (Performance metrics)
    - ♦ Minimize turnaround time
      - » avg time to complete a job
      - »  $T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$
    - ♦ Maximize throughput
      - » operations (jobs) per second
      - » Minimize overhead of context switches: large quanta
      - » Efficient utilization (CPU, memory, disk etc)
    - ♦ Short response time
      - »  $T_{\text{response}} = T_{\text{firstrun}} - T_{\text{arrival}}$
      - » type on a keyboard
      - » Small quanta
    - ♦ Fairness
      - » fair, no starvation, no deadlock
- 

# [lec8] Scheduling policies

---



# [lec8] Multiple Queue Scheduling

---

- Motivation: processes may be of different nature and can be easily classified
  - ♦ e.g. foreground jobs vs. background jobs
- The method:
  - ♦ Processes permanently assigned to one queue, based on processes priority / type
    - » Preference to jobs with higher priorities
  - ♦ Each queue can have its own scheduling algorithm
    - » e.g. RR for foreground queue, FCFS for background queue
  - ♦ Need a scheduling among the queues
    - » e.g. fixed priority preemptive scheduling (high-pri queue trumps other)
    - » e.g. time-slice between queues


# [lec8] Multilevel Feedback Queue (MLFQ)

---

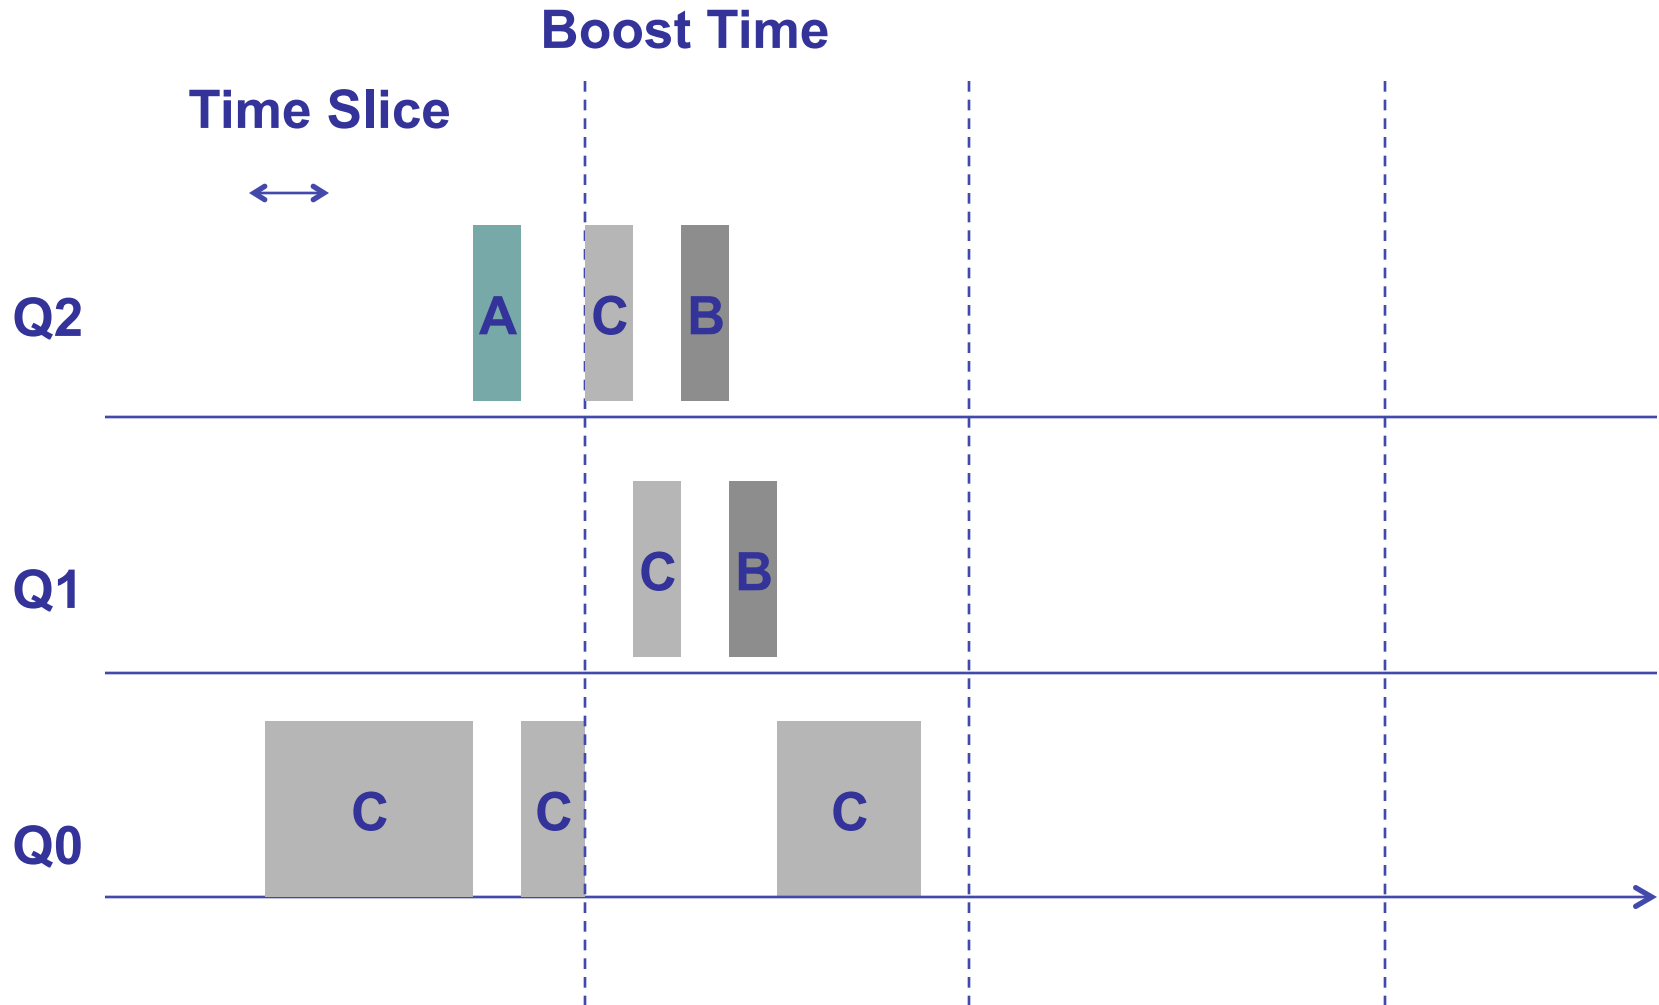
- Problem: how to change priority?
- Jobs start at highest priority queue
- Feedback
  - ♦ If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue).
  - ♦ If a job gives up the CPU before the time slice is up, it stays at the same priority level.
  - ♦ After a long time period, move all the jobs in the system to the topmost queue (aging)



# [lec8] MLFQ Example – a long job + short jobs in between



# [lec8] MLFQ Example – a long job+short jobs, with boost




# [lec8] Scheduling Overhead

---

- Operating systems aim to minimize overhead
  - ♦ Context switching is not doing any useful work and is pure overhead
  - ♦ Overhead includes context switch + making a scheduling decision
- Modern time-sharing OSes (Unix, Windows, ...) time-slice processes in ready list
  - ♦ A process runs for its quantum, OS context switches to another, next process runs, etc.
  - ♦ A CPU-bound process will use its entire quantum (e.g., 10ms)
  - ♦ An IO-bound process will use part (e.g., 1ms), then issue IO
  - ♦ The IO-bound process goes on a wait queue, the OS switches to the next process to run, the IO-bound process goes back on the ready list when the IO completes

# [lec8] CPU Utilization



- CPU utilization is the fraction of time the system is doing useful work (e.g., not context switching)
- If the system has
  - ♦ Quantum of 10ms + context-switch and decision making overhead of 0.1ms
  - ♦ 3 CPU-bound processes + round-robin scheduling
- In steady-state, time is spent as follows:
  - ♦ 10ms + 0.1ms + 10ms + 0.1ms + 10ms + 0.1ms
  - ♦ CPU utilization = time doing useful work / total time
  - ♦ CPU utilization =  $(3 \times 10\text{ms}) / (3 \times 10\text{ms} + 3 \times 0.1\text{ms}) = 30/30.3$
- If one process is IO-bound, it will not use full quantum
  - ♦ 10ms + 0.1ms + 10ms + 0.1ms + 1ms + 0.1ms
  - ♦ CPU util =  $(2 \times 10 + 1) / (2 \times 10 + 1 + 3 \times 0.1) = 21/21.3$

# [lec8] Scheduling Summary

---

- Scheduler (dispatcher) is the module that gets invoked when a context switch needs to happen
- Scheduling algorithm determines which process runs, where processes are placed on queues
- Many potential goals of scheduling algorithms
  - ♦ Utilization, throughput, wait time, response time, etc.
- Various algorithms to meet these goals
  - ♦ FCFS/FIFO, SJF, Priority, RR
- Can combine algorithms
  - ♦ Multiple-level feedback queues

# Memory Management

---

Next few lectures are going to cover memory management

- Goals of memory management
  - ♦ To provide a convenient abstraction for programming
  - ♦ To allocate scarce memory resources among competing processes to maximize performance with minimal overhead
- Mechanisms
  - ♦ Physical and virtual addressing
  - ♦ Techniques: partitioning, paging, segmentation
  - ♦ Page table management, TLBs, VM tricks
- Policies
  - ♦ Page replacement algorithms

# Virtual Memory

---

- The abstraction that the OS provides for managing memory is **virtual memory** (VM)
  - ♦ Virtual memory enables a program to execute with less than its complete data in physical memory
    - » A program can run on a machine with less memory than it “needs”
    - » Can also run on a machine with “too much” physical memory
  - ♦ Many programs do not need all of their code and data at once (or ever) – no need to allocate memory for it
  - ♦ OS will adjust amount of memory allocated to a process based upon its behavior
  - ♦ VM requires hardware support and OS management algorithms to pull it off
- Let’s go back to the beginning...

# In the beginning...

---


- Rewind to the very old days
  - ♦ Programs use **physical addresses** directly
  - ♦ OS loads job, runs it, unloads it



# 1. Simple uniprogramming: Single segment per process

---

Physical memory



# Simple uniprogramming:

## Single segment per process

---

- Highest memory holds OS
- Process is allocated memory starting at 0, up to the OS area
- The single segment contains code, data, stack, heap
- When loading a process, just bring it in at 0
  - ♦ Directly using physical addresses
- Examples:
  - ♦ early batch monitor which ran only one job at a time
    - » if the job wrecks the OS, reboot OS
  - ♦ 1<sup>st</sup> generation PCs operated in a similar fashion
- Pros / Cons?

# Multiprogramming

---

- Want to let several processes coexist in main memory

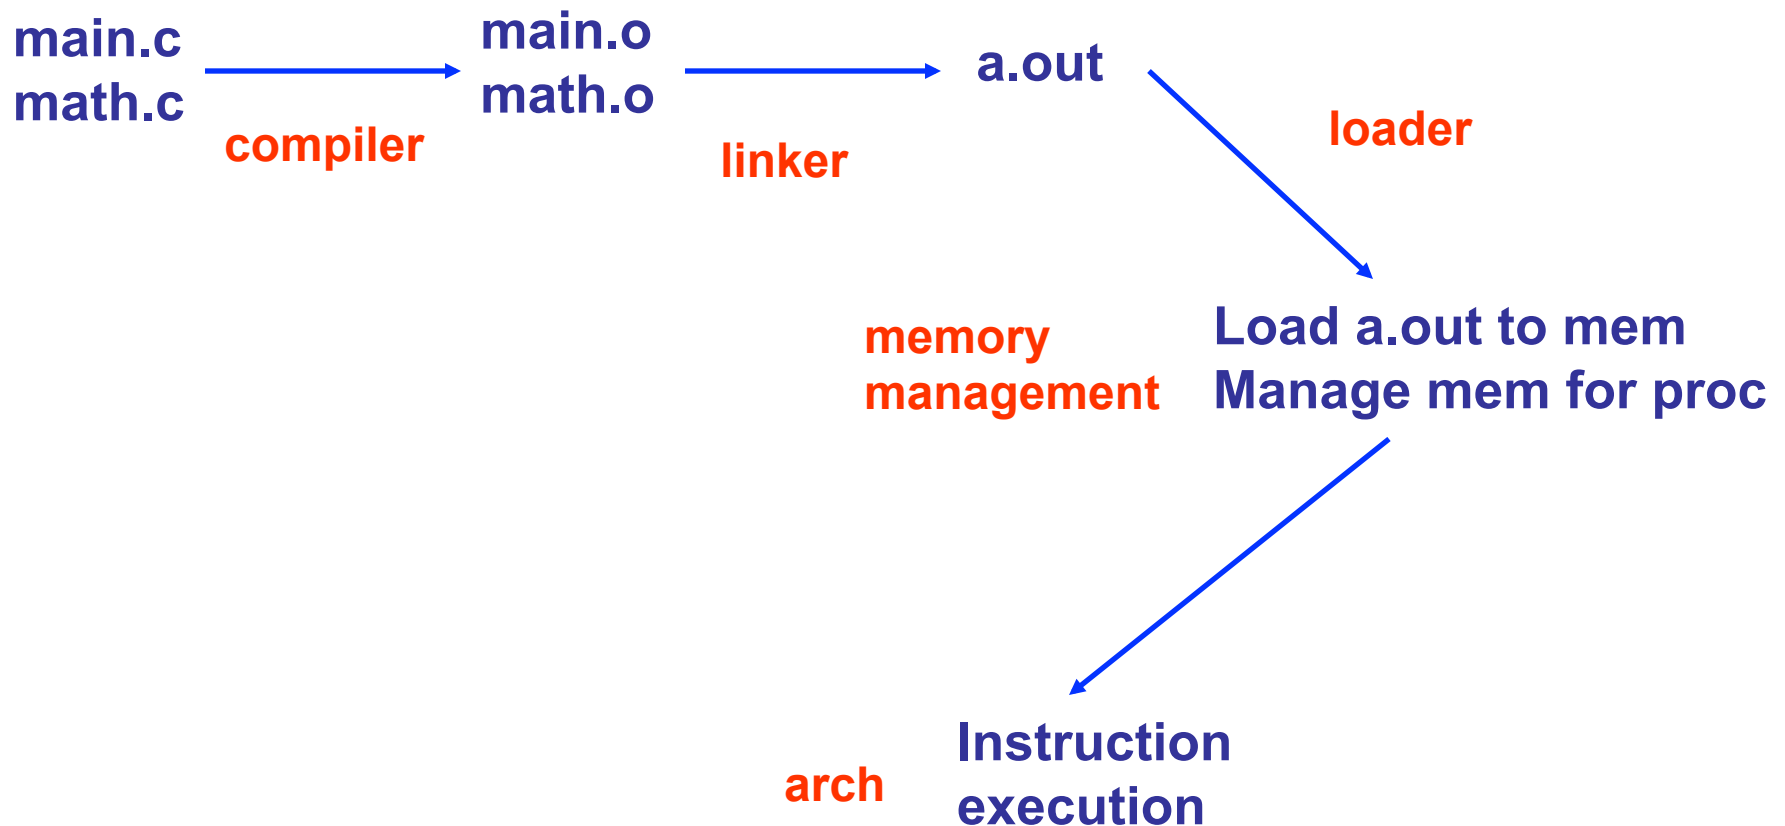
# Issues in sharing main memory

---

- **Transparency:**
  - ◆ Processes should not know memory is shared
  - ◆ Run regardless of the number/locations of processes
- **Safety:**
  - ◆ Processes cannot corrupt each other
- **Efficiency:**
  - ◆ Both CPU and memory utilization shouldn't be degraded badly by sharing

# The Big Picture

---



## 2. Simple multiprogramming

---

With **static software memory relocation**, no protection, 1 segment per process:

- Highest memory holds OS
- Processes allocated memory starting at 0, up to the OS area
- When a process is loaded, **relocate** it so that it can run in its allocated memory area

# Simple multiprogramming:

## Single segment per process, static relocation


---



# Simple multiprogramming:

## Single segment per process, static relocation

---





# Simple multiprogramming:

## Single segment per process, static relocation

---

- four drawbacks
  1. No protection
  2. Low utilization -- Cannot relocate dynamically
    - » Addresses in binary is fixed (after loading)
    - » Cannot do anything about holes
  3. No sharing -- Single segment per process
    - » Cannot share part of process address space (e.g. text)
  4. Entire address space needs to fit in mem
    - » Need to swap whole, very expensive!

# What else can we do?

---

- Already tried
  - ♦ Compile time / linking time
  - ♦ Loading time
- Let us try execution time!

# 3. Dynamic memory relocation

---

- Instead of changing the address of a program before it's loaded, change the address dynamically *during every reference*

*Can this be done in software?*

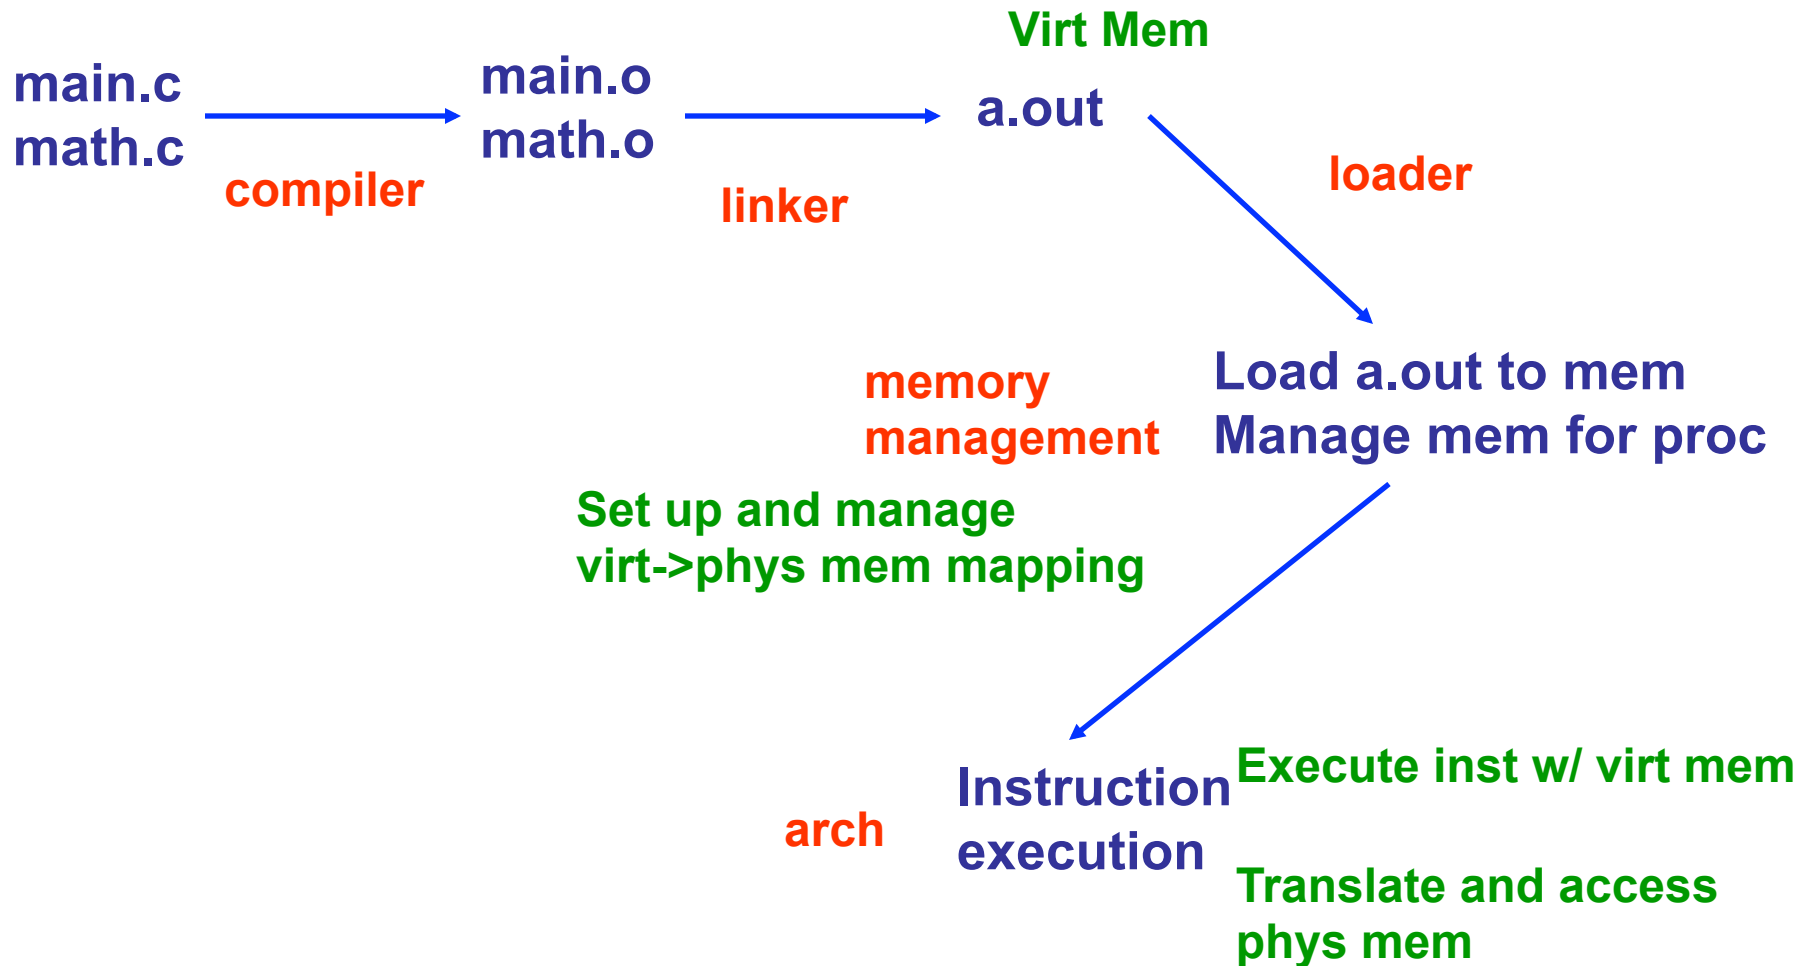
# Virtual Addresses

---

- To make it easier to manage the memory of processes running in the system, we're going to make them use **virtual addresses** (logical addresses)
  - ♦ Virtual addresses are independent of the actual physical location of the data referenced
  - ♦ OS determines location of data in physical memory
  - ♦ Compiler+linker determines virtual memory. OS also allocates virtual memory (heap memory)
  - ♦ CPU executes instructions with virtual addresses
  - ♦ Virtual addresses are translated by hardware into physical addresses (with help from OS)
- The set of virtual addresses that can be used by a process comprises its **virtual address space (VAS)**
  - ♦ VAS often larger than physical memory (64-bit addresses)
  - ♦ But can also be smaller (32-bit VAS with 8 GB of memory)


# The Big Picture

---



# Translation overview


---



- Actual translation process is usually performed by hardware
- Translation table is set up by software
- CPU view
  - ♦ what program sees, virtual addresses
- Memory view
  - ♦ physical memory addresses


# 3.1 Base and bound

---



- Built in Cray-1 (1976)
- A program can only access physical memory in  $[base, base+bound]$
- On a context switch: save/restore base, bound registers
- Pros:
  - ♦ simple, fast translation, cheap
  - ♦ Can relocate segment at execution time


# 3.1 Base and bound



- The essence:
  - ♦ A level of indirection
  - ♦  $\text{Phy. Addr} = \text{Vir. Addr} + \text{base}$
- Why do we need the limit register? Protection
  - ♦ If (physical address > base + limit) then an exception will happen



# 3.1 Base and bound




- Cons:
  - ♦ Relocation requires moving the entire address space
  - ♦ Only one segment per process
  - ♦ How can two processes share code while keeping private data areas?
    - » Can it be done safely with a single-segment scheme?

# What have we solved?

---

- four drawbacks
  1. No protection
  2. Low utilization -- Cannot relocate dynamically
    - » Cannot do anything about holes
  3. No sharing -- Single segment per process
    - » Cannot share part of process address space (e.g. text)
  4. Entire address space needs to fit in mem
    - » Need to swap whole, very expensive!

## 3.2 Multiple Segments



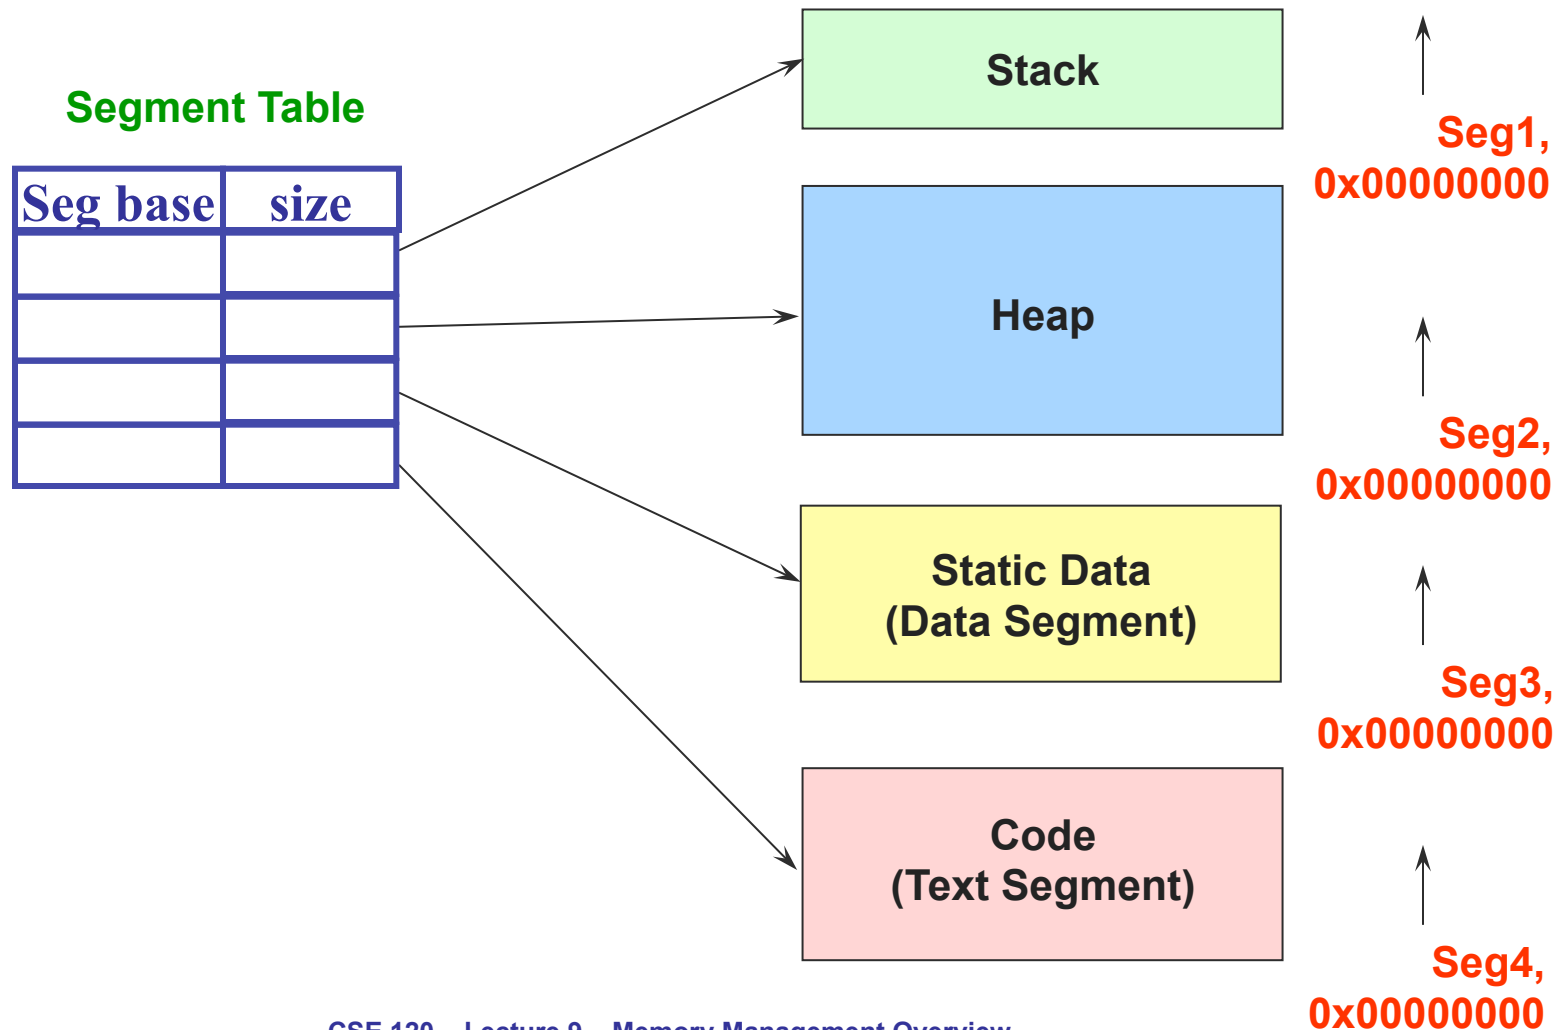
- Separate a virtual memory address space into multiple “segments”
- A hardware segment table of (seg base, size), each entry also has an associated permission (nil, read, write, exec)
- On a context switch: save/restore the table (or a pointer to the table) in kernel memory

# Segmentation

---

- Segmentation is a technique that partitions memory into logically related data units
  - ♦ Module, procedure, stack, data, file, etc.
- Natural extension of base-and-bound
  - ♦ Base-and-bound: 1 segment/process
  - ♦ Segmentation: many segments/process

# Segmented Address Space



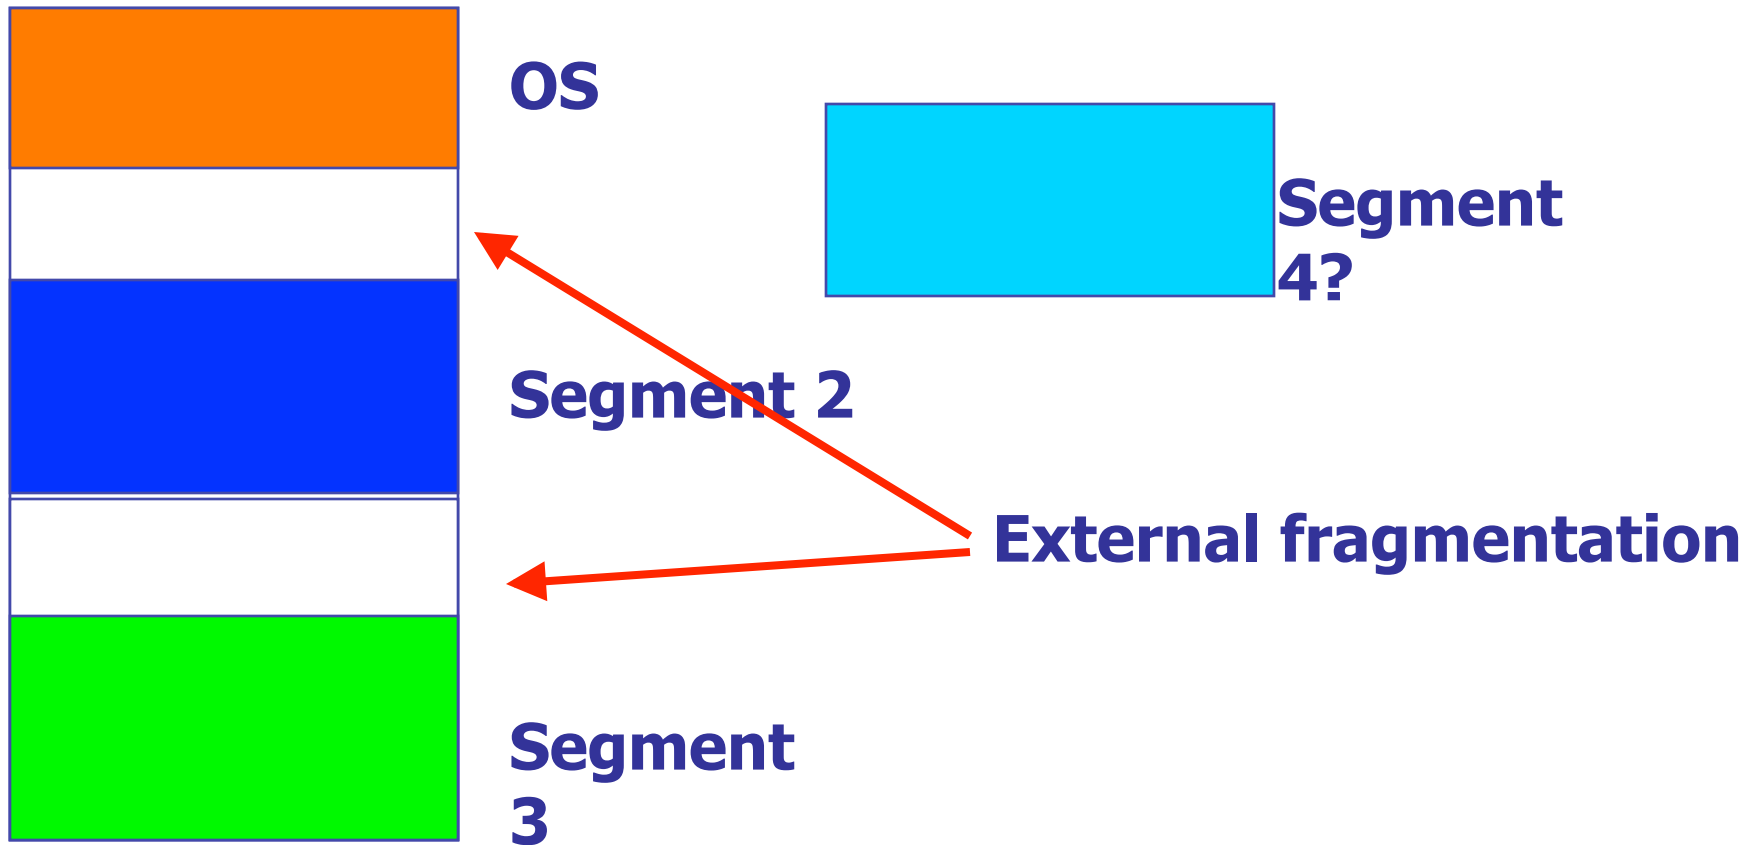
# Pros/cons of segmentation

---

- Pros:
  - ♦ Process can be split among several segments
    - » Allows sharing
  - ♦ Segments can be assigned, moved, or swapped independently
- Cons:
  - ♦ External fragmentation: many holes in physical memory
    - » Also happens in base and bound scheme

# External fragmentation with segmentation

---



# What fundamentally causes external fragmentation?

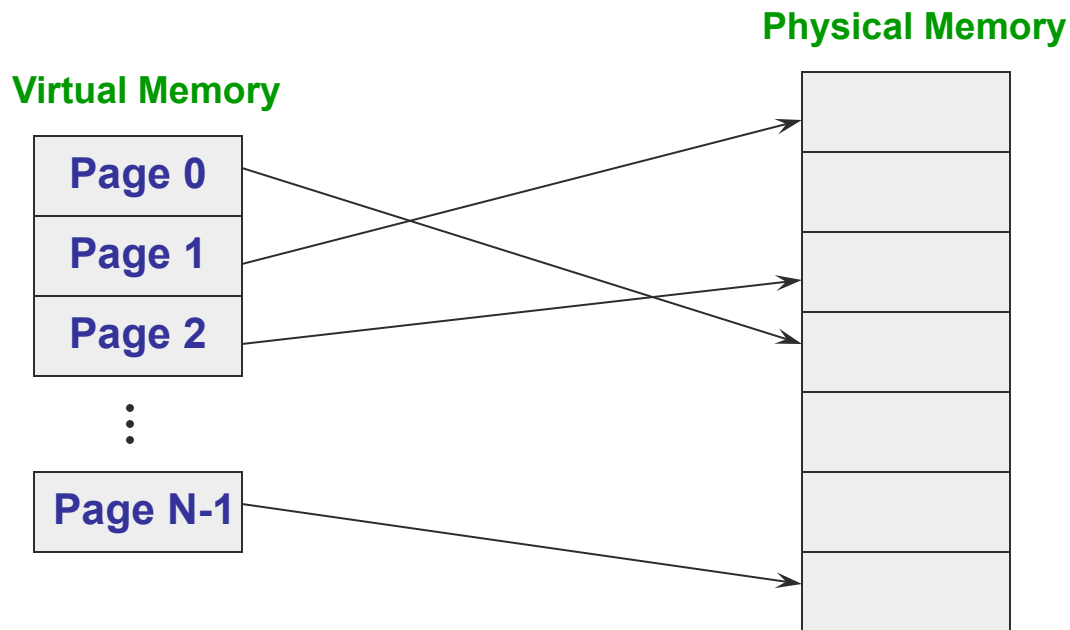
---

1. Segments of many different sizes
  2. Each has to be allocated contiguously
- “Million-dollar” question:  
*Physical memory is precious.*  
*Can we limit the waste to a single hole of  $X$  bytes?*



# Paging

- Paging solves the external fragmentation problem by using fixed sized units in both physical and virtual memory




# Paging

---

- Translating addresses
  - ♦ Virtual address has two parts: **virtual page number** and **offset**
  - ♦ Virtual page number (VPN) is an index into a page table
  - ♦ Page table determines page frame number (PFN)
  - ♦ Physical address is PFN::offset (“::” means concatenate)
- Page tables
  - ♦ Map **virtual page number** (VPN) to **page frame number** (PFN)
    - » VPN is the index into the table that determines PFN
  - ♦ One page table entry (PTE) per page in virtual address space
    - » Or, one PTE per VPN

# Paging



- Context switch
  - ♦ similar to the segmentation scheme
- Pros:
  - ♦ easy to allocate memory
  - ♦ easy to swap
  - ♦ easy to share


# Paging Example

---

- Pages are 4K
  - ♦ 4K  $\rightarrow$  offset is 12 bits  $\rightarrow$  VPN is 20 bits ( $2^{20}$  VPNs), assuming 32bit system
- Virtual address is 0x7468
  - ♦ Virtual page is 0x7, offset is 0x468 (lowest 12 bits of address)
- Page table entry 0x7 contains 0x2
  - ♦ Page frame number is 0x2
  - ♦ Seventh virtual page is at address 0x2000 (physical page 2)
- Physical address = 0x2000 :: 0x468 = 0x2468

# Deep thinking: Paging implementation

- Translation: table lookup and bit substitution
- Why is this possible?
- Why can't we do the same in segmentation?



# Summary

---

- Virtual memory
  - ♦ Processes use virtual addresses
  - ♦ Hardware translates virtual address into physical addresses with OS support
- Evolution of techniques
  - ♦ Single, fixed physical segment per process (no virt mem)
  - ♦ Single segment per process, static relocation (no virt mem)
  - ♦ Base-and-bound – dynamic relocating whole process
  - ♦ Segmentation – multiple (variable-size) segments with dynamic relocation
  - ♦ Paging – small, fixed size pages

# Next time...

---

- Chapters 18, 19, 20