

CSE 120

Principles of Operating Systems

Fall 2021

Lecture 10: Paging


Yiying Zhang

Lecture Overview


We'll cover more paging mechanisms:

- Optimizations
 - ♦ Managing page tables (space)
 - ♦ Efficient translations (TLBs) (time)
 - ♦ Demand paged virtual memory (space), next lecture
- Midterm grades to be released soon
- Homework 3 out
- Work on your project 2!

**All problems in
computer
science can be
solved by
another level of
indirection**



All problems in computer science can be solved by another level of indirection



Butler Lampson

David Wheeler

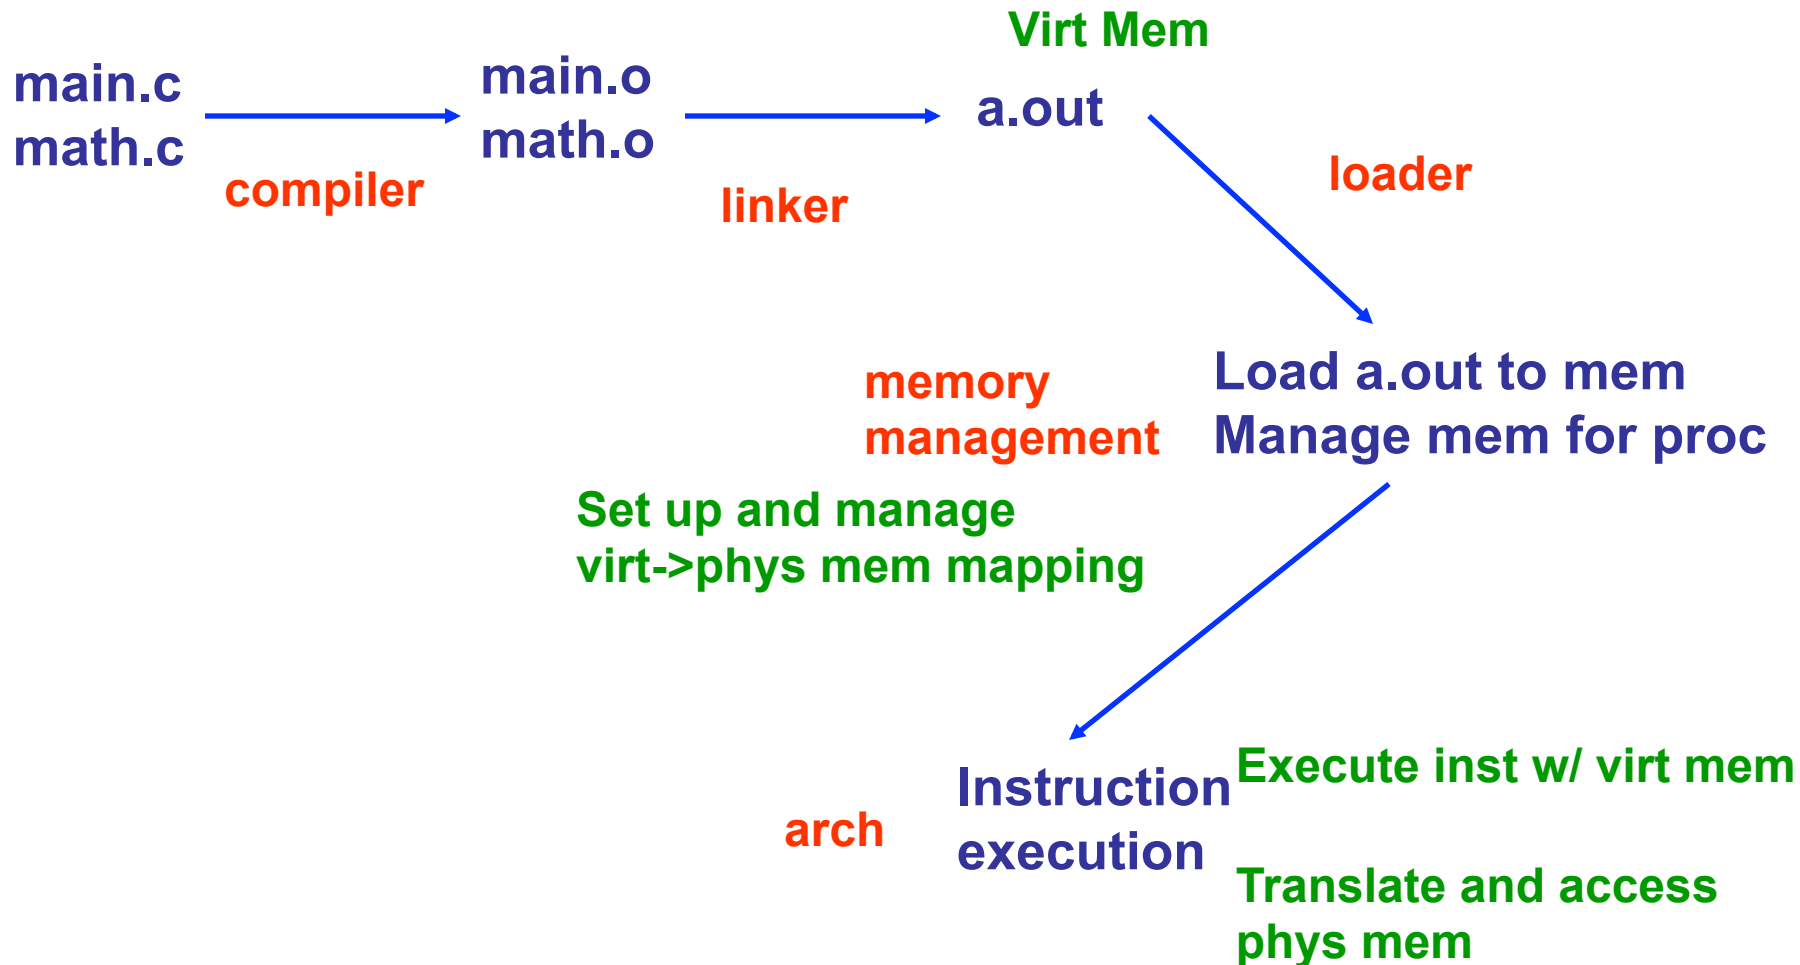
but that usually will create another problem

– David Wheeler

Summary: Evolution of Memory Management (before paging)

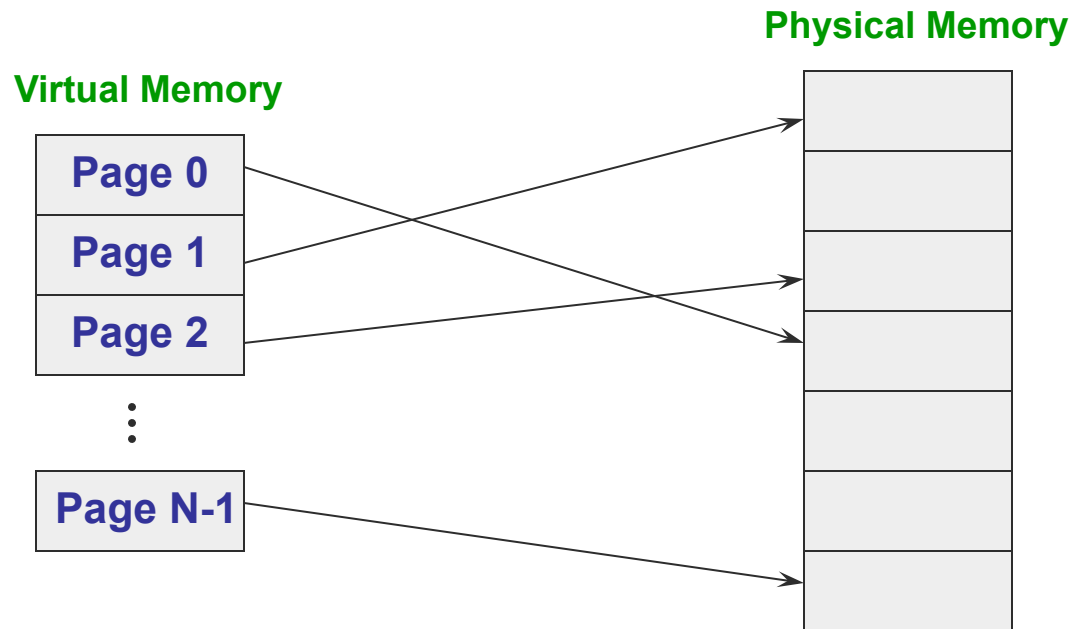
Scheme	How	Pros	Cons
Simple uniprogramming	1 segment loaded to starting address 0	Simple	1 process 1 segment No protection
Simple multiprogramming	1 segment relocated at loading time	Simple, Multiple processes	1 segment/process No protection External frag.
Base & Bound	Dynamic mem relocation at runtime	Simple hardware, Multiple processes Protection	1 segment/process, External frag.
Multiple segments	Dynamic mem relocation at runtime	Sharing, Protection, multi segs/process	More hardware, External frag.

[lec9] The Big Picture



[lec9] Paging


- Paging solves the external fragmentation problem by using fixed sized units in both physical and virtual memory



[lec9] Paging

- Translating addresses
 - ♦ Virtual address has two parts: **virtual page number** and **offset**
 - ♦ Virtual page number (VPN) is an index into a page table
 - ♦ Page table determines page frame number (PFN)
 - ♦ Physical address is PFN::offset (“::” means concatenate)
- Page tables
 - ♦ Map **virtual page number** (VPN) to **page frame number** (PFN)
 - » VPN is the index into the table that determines PFN
 - ♦ One page table entry (PTE) per page in virtual address space
 - » Or, one PTE per VPN

[lec9] Paging



- Context switch
 - ♦ similar to the segmentation scheme
- Pros:
 - ♦ easy to allocate memory
 - ♦ easy to swap
 - ♦ easy to share

[lec9] Paging Example

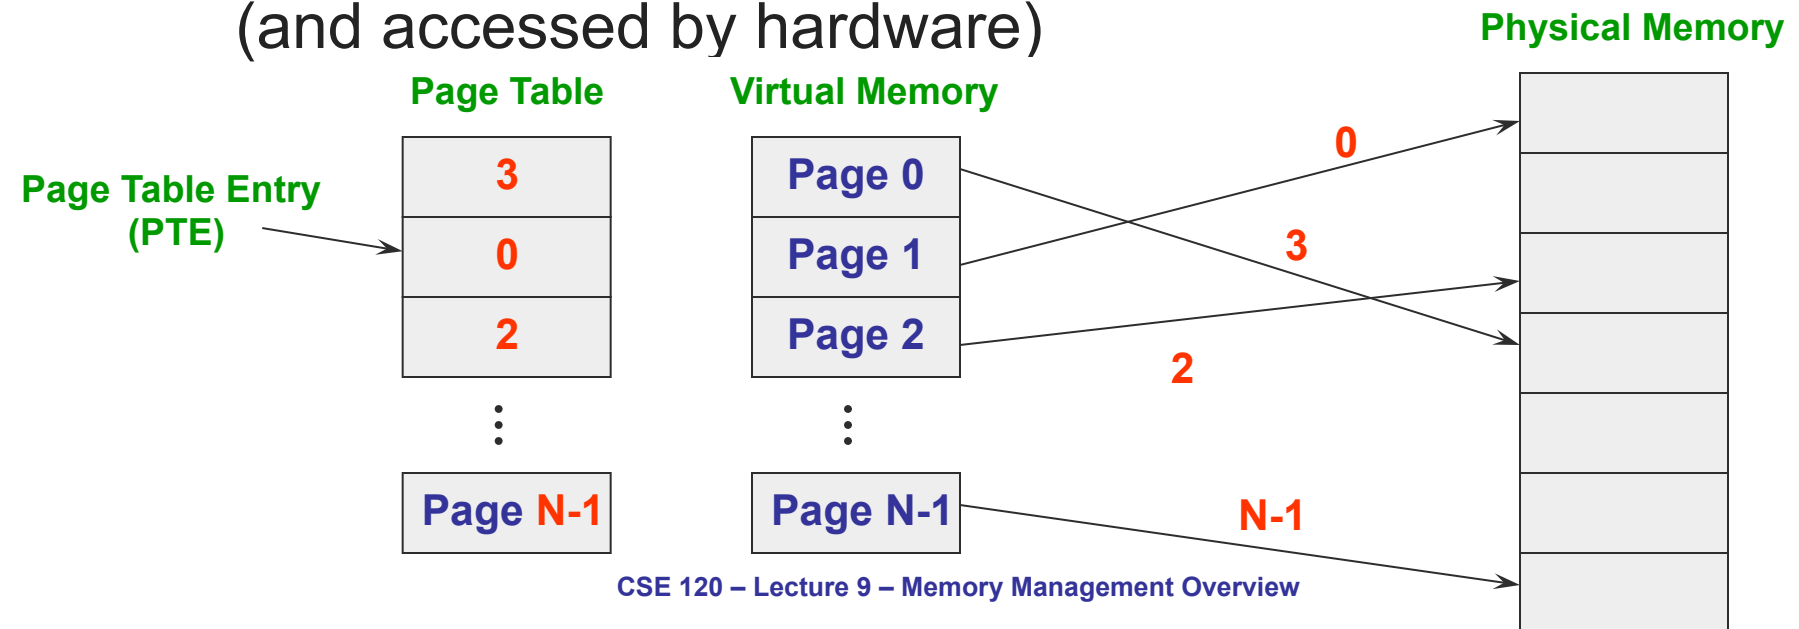
- Pages are 4K
 - ♦ 4K \rightarrow offset is 12 bits \rightarrow VPN is 20 bits (2^{20} VPNs), assuming 32bit system
- Virtual address is 0x7468
 - ♦ Virtual page is 0x7, offset is 0x468 (lowest 12 bits of address)
- Page table entry 0x7 contains 0x2
 - ♦ Page frame number is 0x2
 - ♦ Seventh virtual page is at address 0x2000 (physical page 2)
- Physical address = 0x2000 :: 0x468 = 0x2468

Summary so far

- Virtual memory
 - ♦ Processes use virtual addresses
 - ♦ Hardware translates virtual address into physical addresses with OS support
- Evolution of techniques
 - ♦ Single, fixed physical segment per process (no virt mem)
 - ♦ Single segment per process, static relocation (no virt mem)
 - ♦ Base-and-bound – dynamic relocating whole process
 - ♦ Segmentation – multiple (variable-size) segments with dynamic relocation
 - ♦ Paging – small, fixed size pages

Page Tables

- Page tables completely define the mapping between virtual pages and physical pages for an address space
- Each process has an address space, so each process has a page table
- Page tables are data structures maintained by the OS (and accessed by hardware)



Page Table Entries (PTEs)

1	1	1	3	20
M	R	V	Prot	Page Frame Number

- Page table entries control mapping
 - ♦ The **Modify** bit says whether or not the page has been written
 - » It is set when a write to the page occurs
 - ♦ The **Reference** bit says whether the page has been accessed
 - » It is set when a read or write to the page occurs
 - ♦ The **Valid** bit says whether or not the PTE can be used
 - » It is checked each time the virtual address is used
 - ♦ The **Protection** bits say what operations are allowed on page
 - » Read, write, execute
 - ♦ The **page frame number** (PFN) determines physical page

Paging implementation – how does it really work?

- Where to store page table?
- How to use MMU?
 - ♦ Even small page tables are too large to load into MMU
 - ♦ Page tables kept in mem and MMU only has their base addresses
- What happens at context switches?

Paging Advantages

- Easy to allocate (physical) memory
 - ♦ Memory comes from a free list of fix-sized chunks
 - ♦ Allocating a page is just removing it from the list
 - ♦ External fragmentation not a problem
- Easy to swap out chunks of a program
 - ♦ All chunks are the same size
 - ♦ Use valid bit to detect references to swapped pages
 - ♦ Pages are a convenient multiple of the disk block size
 - ♦ More on swapping next time

Paging Limitations

- Can still have **internal fragmentation**
 - ♦ Process may not use memory in multiples of a page
- Memory reference overhead
 - ♦ 2 references per address lookup (page table, then memory)
 - ♦ **Solution – use a hardware cache of lookups (next lec)**
- Memory required to hold page table can be significant

Deep thinking

- Why does the page table we talked about so far have to be contiguous in the physical memory?
 - ♦ Why did a segment have to be contiguous in memory?
- For a 4GB virtual address space, we just need 1M PTE (~4MB), what is the big deal?
- My PC has 2GB, why do we need PTEs for the entire 4GB address space?

How many PTEs do we need?

(assume page size is 4096 bytes)


- Worst case for 32-bit address machine?
- What about 64-bit address machine?
- Page size?
 - ♦ Small page -> big table
 - ♦ Large page -> small table but large internal fragmentation

Managing Page Tables


- How can we reduce page table space overhead?
 - ♦ Observation: Only need to map the portion of the address space actually being used (tiny fraction of entire addr space)
- How can we be flexible?
 - “All computer science problems can be solved with an extra level of indirection.”
 - two-level page tables

Page Table Evolution


Linear (Flat)
Page Table



Virtual Address
Space




Physical Memory




Page Table Evolution


Hierarchical
Page Table



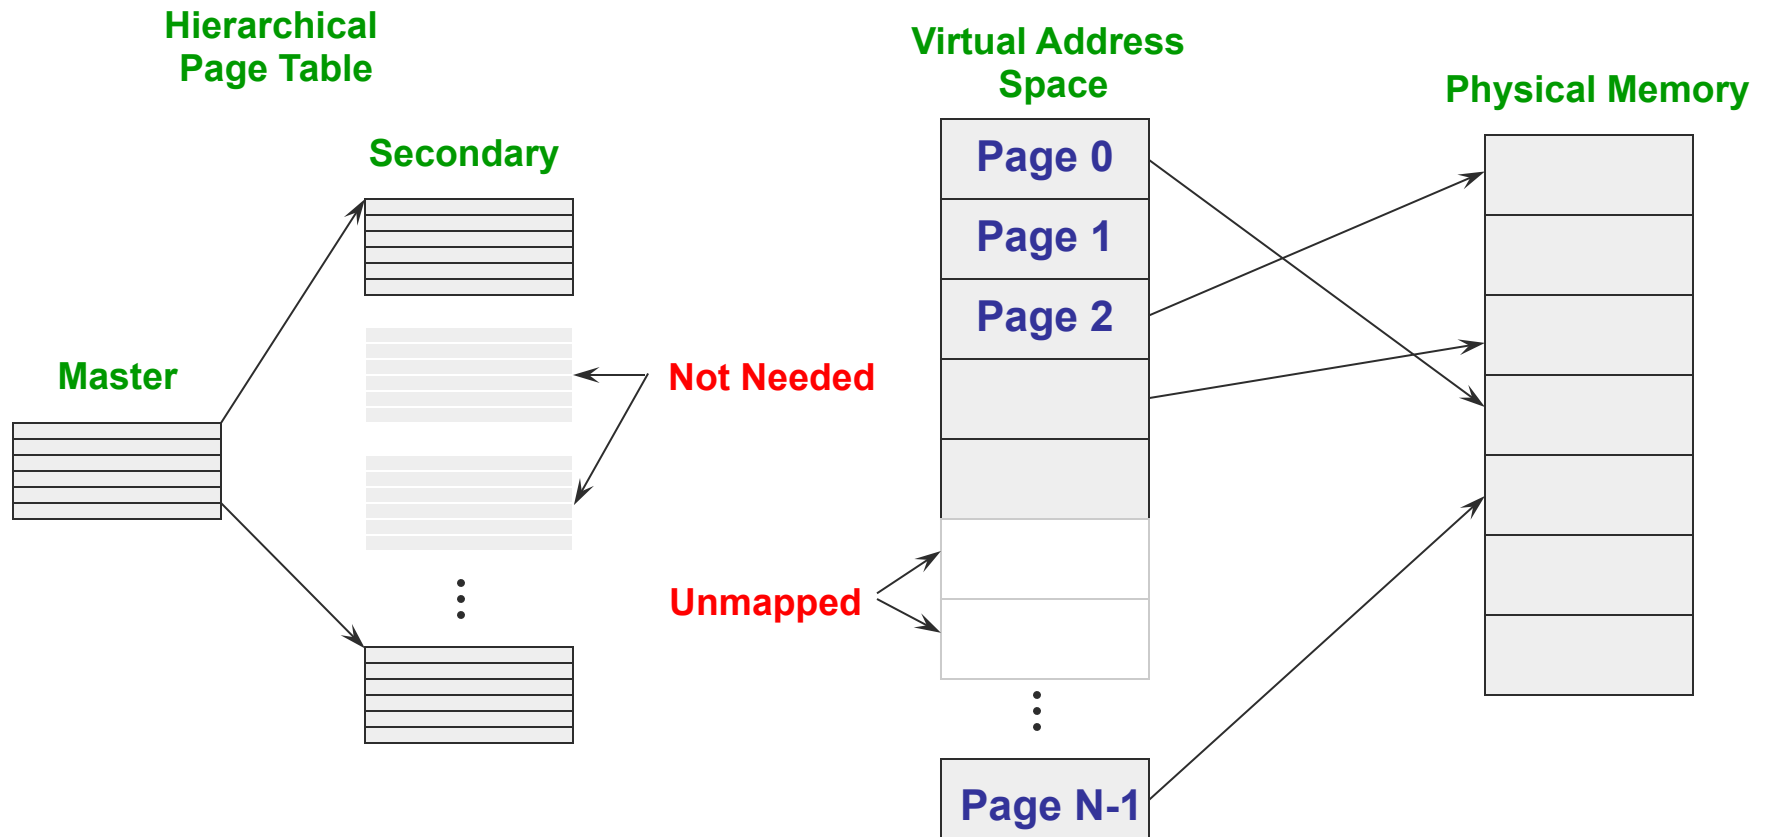
Virtual Address
Space



Physical Memory



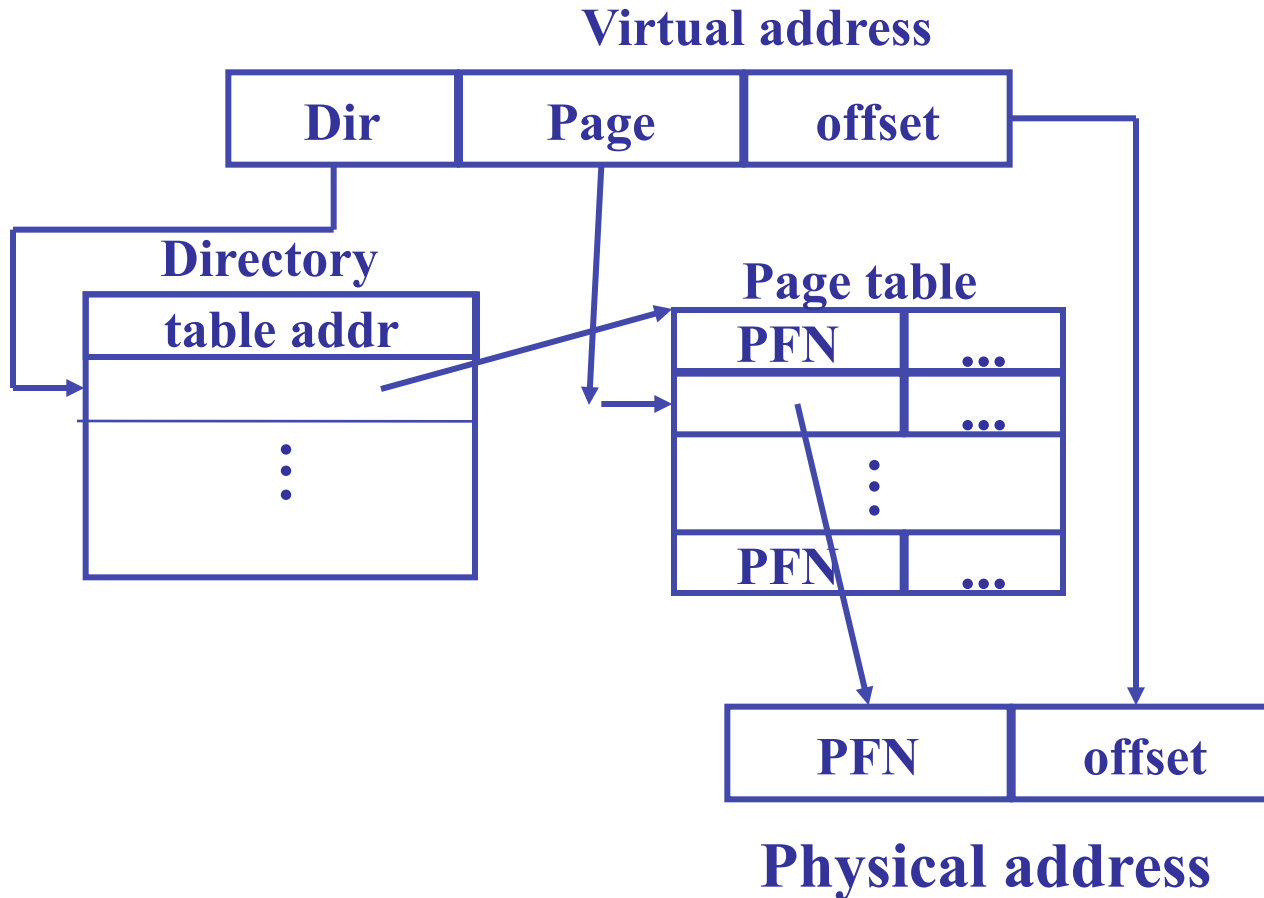
Page Table Evolution



Two-Level Page Tables

- Two-level page tables
 - ♦ Virtual addresses (VAs) have three parts:
 - » Directory (master page number), secondary page number, and offset
 - ♦ Directory page table maps VAs to secondary page table
 - ♦ Secondary page table maps page number to physical page
 - ♦ Offset indicates where in physical page address is located

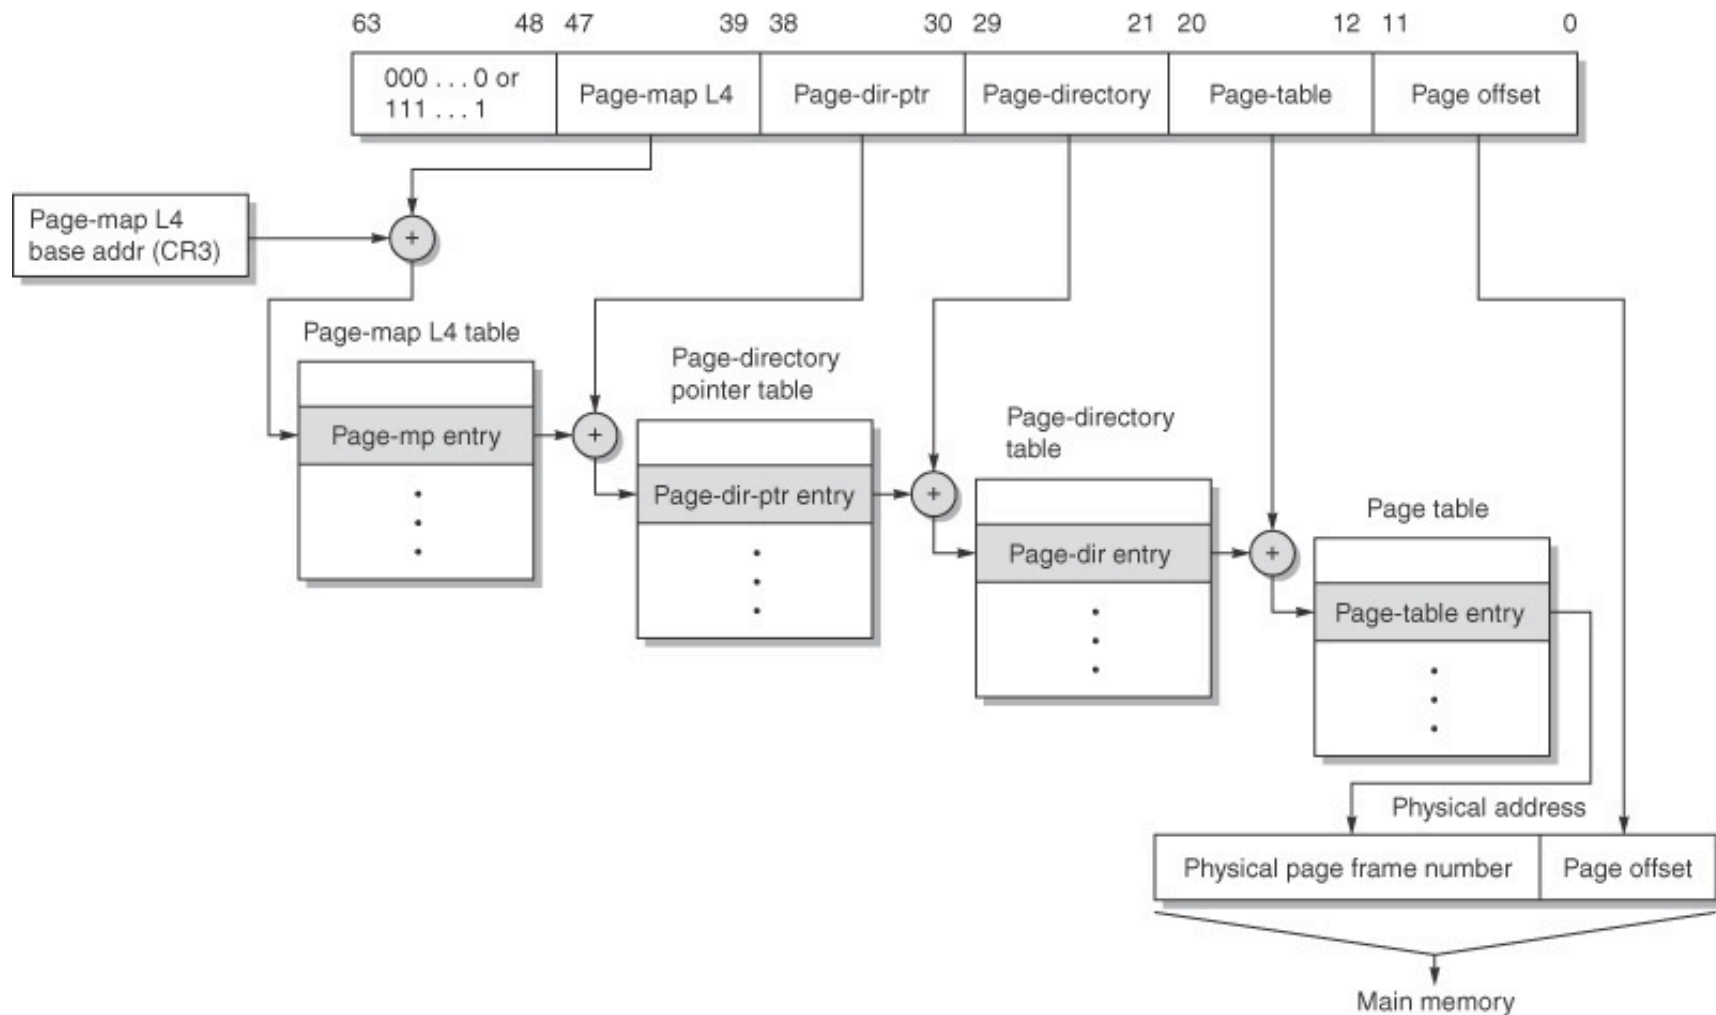
Two-Level Page Tables



Two-Level Page Tables

- Example
 - ♦ 4KB pages, 4 bytes/PTE, 32-bit address space
 - ♦ How many bits in offset?
 - » 4KB = 12 bits
 - ♦ Want directory page table in one page, how many entries can we have in the directory page table?
 - » 4KB/4 bytes = 1KB entries (each entry is a 32-bit address)
 - ♦ Hence, 1KB secondary page tables. How many bits?
 - ♦ Directory (1KB) = 10, offset = 12, $32 - 10 - 12 = 10$ bits left
 - » One secondary page table can host 4K/4bytes=1KB PTEs
 - » 10 bits (inner) => exactly 1KB PTEs

Multiple-level page tables



Multi-level page tables

- 3 Advantages?
 - ♦ L1, L2, L3 tables do not have to be consecutive
 - ♦ They do not have to be allocated before use!
 - ♦ They can be swapped out to disk!

The power of an extra level of indirection!

- Problems?

Efficient Translations

- Our original page table scheme already increased the cost of doing memory lookups
 - ♦ Two lookups into the page table, another to fetch the data
 - ♦ One lookup and one data access for original flat page table
- Now 4-level page tables require five DRAM accesses for one memory operation!
 - ♦ Four lookups into the page tables, a fifth to fetch the data
- Solution: *reference locality!*
 - ♦ In a short period of time, a process is likely accessing only a few pages
 - ♦ Store part of the page table that is “hot” in a fast hardware unit

Next time


- Swapping, memory allocation, memory sharing

Backup Slides

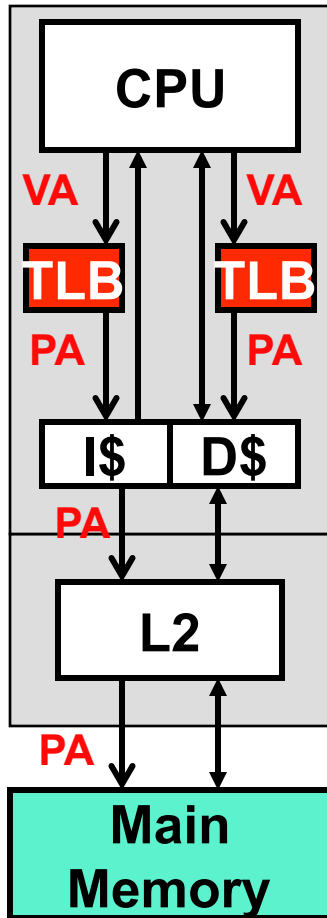
Memory Hierarchy Revisited

What does this
imply about L1
addresses?

Where do we hope
requests get satisfied?

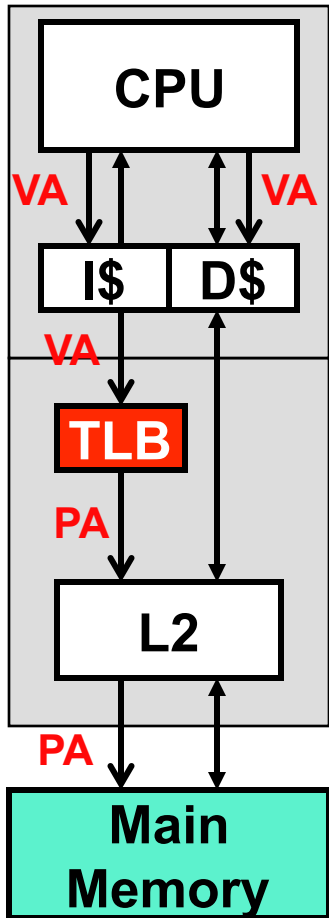


Physical (Address) Caches



- Memory hierarchy so far: **physical caches**
 - ♦ Indexed and tagged by PAs
 - » Physically Indexed (PI)
 - » Physically Tagged (PT)
 - ♦ Translate to PA to VA at the outset
 - + Cached inter-process communication works
 - » Single copy indexed by PA
 - Slow: adds at least one cycle to t_{hit}

Virtual Caches (VI/VT)




- Alternative: **virtual caches**
 - ♦ Indexed and tagged by VAs (VI and VT)
 - ♦ Translate to PAs only to access L2
 - + Fast: avoids translation latency in common case
 - Problem: VAs from **different processes** are distinct physical locations (with different values) (call **homonyms**)
- What to do on process switches?
 - ♦ Flush caches? Slow
 - ♦ Add process IDs to cache tags
- Does inter-process communication work?
 - ♦ **Synonyms**: multiple VAs map to same PA
 - » Can't allow same PA in the cache twice
 - » Also a problem for DMA I/O
 - ♦ Can be handled, but very complicated


Memory Hierarchy Re- Revisited

What does this
imply about L1
addresses?

Any speed benefits?
Any drawbacks?



Parallel TLB/Cache Access (VI/PT)



- Compromise: **access TLB in parallel**
 - ♦ *In small caches, index of VA and PA the same*
 - » $VI == PI$
 - ♦ Use the VA to index the cache
 - ♦ Tagged by PAs
 - ♦ Cache access and address translation in parallel
 - + No context-switching/aliasing problems
 - + Fast: no additional t_{hit} cycles
 - ♦ Common organization in processors today