

CSE 120

Principles of Operating Systems

Fall 2021

Lecture 11: TLB, Swapping

Yiying Zhang

Announcements

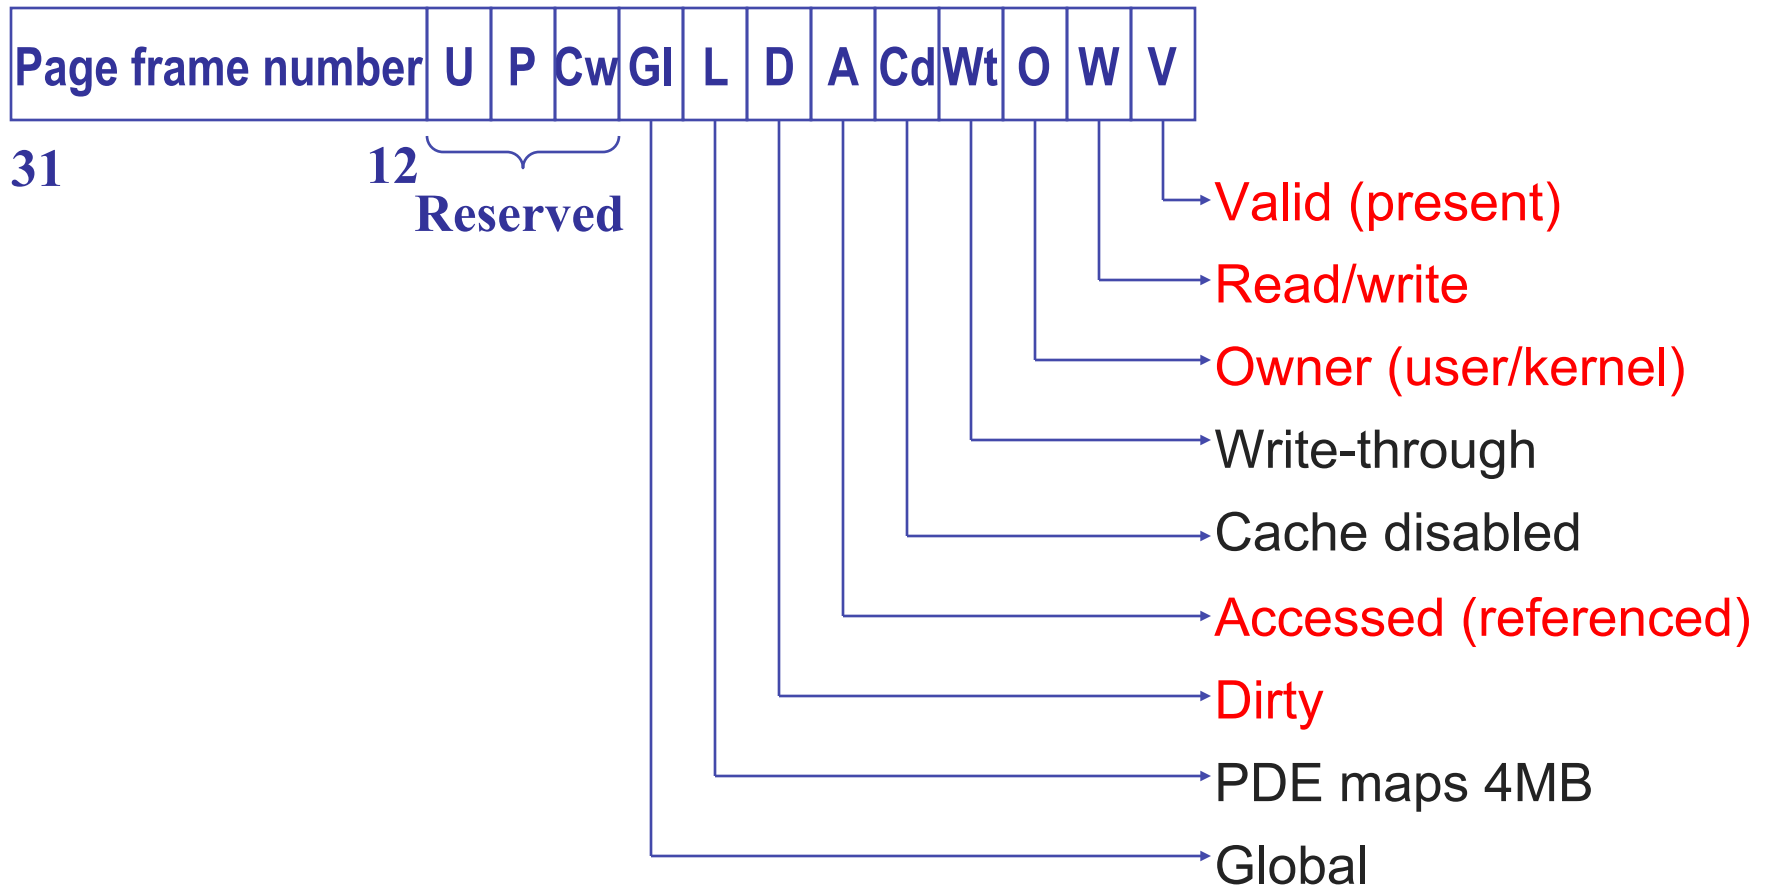
- Midterm graded
- Tomorrow's discussion section will go over PR2 and some of HW3 (if have time)

[lec10] Page Table Entries (PTEs)

1	1	1	3	20
M	R	V	Prot	Page Frame Number

- Page table entries control mapping
 - ♦ The **Modify** bit says whether or not the page has been written
 - » It is set when a write to the page occurs
 - ♦ The **Reference** bit says whether the page has been accessed
 - » It is set when a read or write to the page occurs
 - ♦ The **Valid** bit says whether or not the PTE can be used
 - » It is checked each time the virtual address is used
 - ♦ The **Protection** bits say what operations are allowed on page
 - » Read, write, execute
 - ♦ The **page frame number** (PFN) determines physical page

x86 Page Table Entry



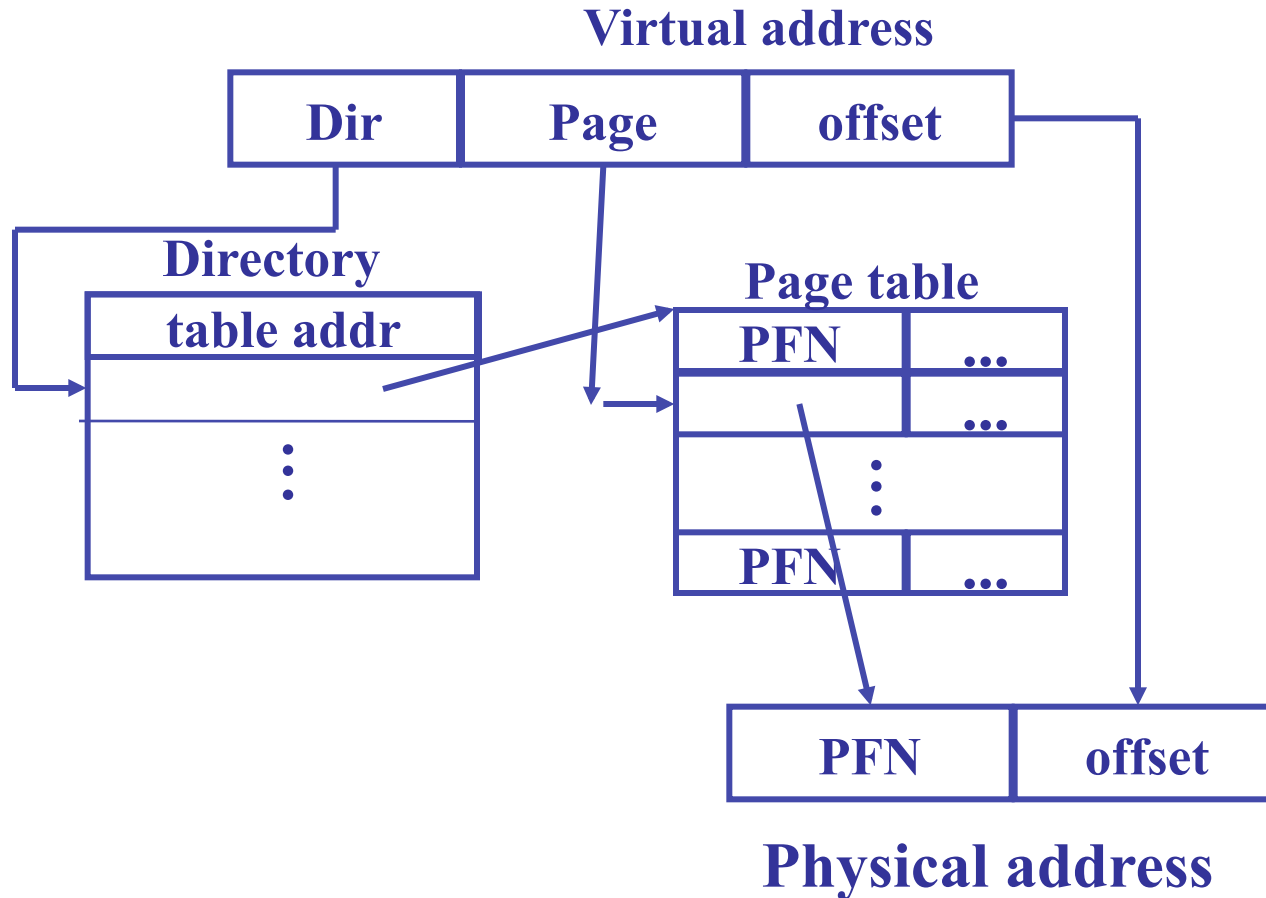
[lec10] Paging implementation – how does it really work?

- Where to store page table?
- How to use MMU?
 - ♦ Even small page tables are too large to load into MMU
 - ♦ Page tables kept in mem and MMU only has their base addresses
- What happens at context switches?


[lec10] Managing Page Tables

- How can we reduce page table space overhead?
 - ♦ Observation: Only need to map the portion of the address space actually being used (tiny fraction of entire addr space)
- How can we be flexible?
 - “All computer science problems can be solved with an extra level of indirection.”
 - two-level page tables

[lec10] Two-Level Page Tables



[lec10] Multiple-level page tables



[lec10] Multi-level page tables

- 3 Advantages?
 - ♦ L1, L2, L3 tables do not have to be consecutive
 - ♦ They do not have to be allocated before use!
 - ♦ They can be swapped out to disk!

The power of an extra level of indirection!

- Problems?

[lec10] Efficient Translations

- Our original page table scheme already increased the cost of doing memory lookups
 - ♦ Two lookups into the page table, another to fetch the data
 - ♦ One lookup and one data access for original flat page table
- Now 4-level page tables require five DRAM accesses for one memory operation!
 - ♦ Four lookups into the page tables, a fifth to fetch the data
- Solution: reference locality!
 - ♦ In a short period of time, a process is likely accessing only a few pages
 - ♦ Store part of the page table that is “hot” in a fast hardware unit

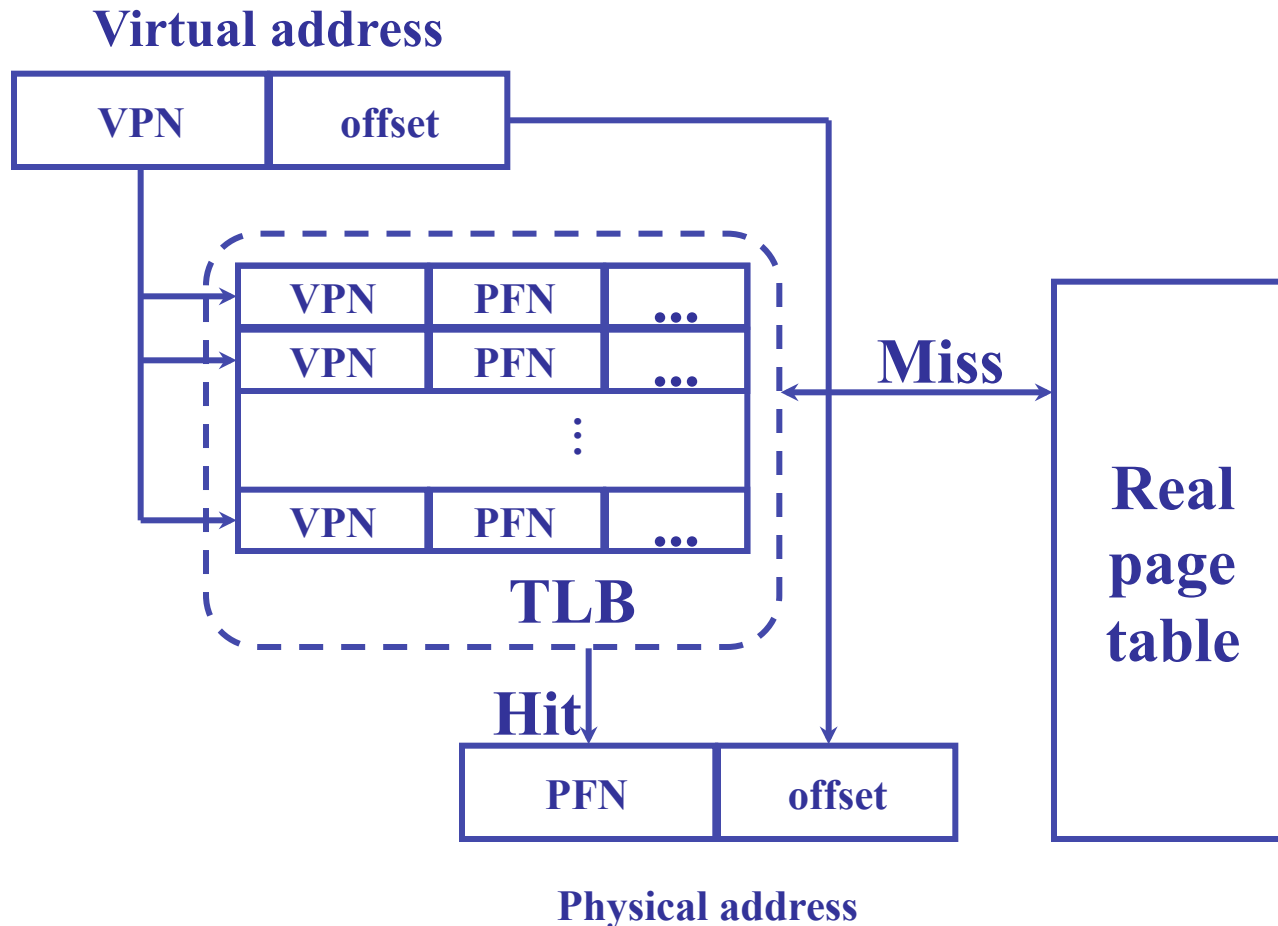
Translation Look-aside Buffer (TLB)

- Translation Look-aside Buffers
 - ♦ Translate VPNs into PFNs
- TLBs implemented in hardware
 - ♦ TLB hit is very fast ≤ 1 CPU cycle
 - ♦ Fully associative cache \Rightarrow least conflict misses
 - ♦ New entries can be inserted anywhere in the TLB
 - ♦ All entries looked up in parallel
 - \Rightarrow TLB can't be made very big, typically 64 – 4096 entries
- Optional (useful) bits
 - ♦ ASIDs -- Address-space identifiers (process tags)

VPN	PFN	...
VPN	PFN	...
⋮		
VPN	PFN	...

TLB

Translation Look-aside Buffer (TLB)



Miss handling:

Hardware-controlled TLB

- On a TLB hit, MMU checks the valid bit
 - ♦ If valid, perform address translation
 - ♦ If invalid (e.g. page not in memory), MMU generates a page fault
 - » OS performs fault handling
 - » Restart the faulting instruction
- On a TLB miss
 - ♦ MMU parses page table and loads PTE into TLB
 - » Needs to replace if TLB is full
 - » Page table layout is fixed
 - ♦ Same as hit ...

Miss handling:

Software-controlled TLB

- On a TLB hit, MMU checks the valid bit
 - ♦ If valid, perform address translation
 - ♦ If invalid (e.g. page not in memory), MMU generates a page fault
 - » OS performs page fault handling
 - » Restart the faulting instruction
- On a TLB miss, HW raises exception, traps to the OS
 - ♦ OS parses page table and loads PTE into TLB
 - » Needs to replace if TLB is full
 - » Page table layout can be flexible
 - ♦ Same as in a hit...

Hardware vs. software controlled

- Hardware approach
 - ♦ Efficient – TLB misses handled by hardware
 - ♦ OS intervention is required only in case of page fault
 - ♦ Page structure prescribed by MMU hardware -- rigid
- Software approach
 - Less efficient -- TLB misses are handled by software
 - MMU hardware very simple, permitting larger, faster TLB
 - OS designer has complete flexibility in choice of MM data structure

Deep thinking

- Without TLB, how MMU finds PTE is fixed
- With TLB, it can be flexible, e.g. software-controlled is possible
- What enables this?
- TLB is an extra level of indirection!

More TLB Issues

- When the TLB misses and a new PTE has to be loaded, a cached PTE must be evicted
 - ♦ Which TLB entry should be replaced?
 - » Random
 - » LRU
- What happens when changing a page table entry (e.g. because of swapping, change read/write permission)?
 - ♦ Change the entry in memory
 - ♦ flush (eg. invalidate) the TLB entry
 - » INGLPG on x86

What happens to TLB in a process context switch?

- During a process context switch, cached translations can not be used by the next process
 - ♦ Invalidate all entries during a context switch
 - » Lots of TLB misses afterwards
 - ♦ Tag each entry with an ASID
 - » Add a HW register that contains the process id of the current executing process
 - » TLB hits if an entry's process id matches that register

Cache vs. TLB

- Similarities:
 - ♦ Both cache a part of the physical memory
- Differences:
 - ♦ Associatively
 - » TLB is usually fully associative
 - » Cache can be direct mapped
 - ♦ Coherence
 - » No hardware provided coherence between TLB and main memory
 - » Software needs to flush TLB entries for coherence
 - » Cache: hardware-provided (via snooping bus) coherence across multiple cores and main memory

More on coherence issues

- No hardware maintains coherence between DRAM and TLBs:
 - ♦ OS needs to flush related TLBs whenever changing a page table entry in memory
- On multiprocessors, when you modify a page table entry, you need to do “*TLB shoot-down*” to flush all related TLB entries at all the cores

Summary so far

- Virtual memory addresses: a level of indirection to decouple static time (compiler) from run time (OS)
- Paging: avoiding external fragmentation, great flexibility
- Single-level page tables are too big
- Multi-level page tables reduce the space overhead (leveraging indirection) but increases the performance overhead
- TLB improves paging performance (leveraging locality)
- But TLB shutdown is costly (esp. on many cores)

Remaining of This Lecture

We'll cover more virtual memory topics:

- Optimizations
 - ♦ Managing page tables (space)
 - ♦ Efficient translations (TLBs) (time)
 - ♦ Demand paged virtual memory (swapping) (space)
- Memory allocation
- Kernel address space (if have time)

[lec9] Sharing main memory

- Simple multiprogramming – 4 drawbacks
 - ♦ Lack of protection
 - ♦ Cannot relocate dynamically
 - dynamic memory relocation: base&bound
 - ♦ Single segment per process
 - dynamic memory relocation: segmentation, paging
 - ♦ Entire address space needs to fit in mem
 - » More need for swapping
 - » Need to swap whole, very expensive!

The last drawback

- So far we've separated the process's view of memory from the OS's view using a mapping mechanism
 - ♦ Each sees a different organization
 - ♦ Allows OS to shuffle processes around
 - ♦ Simplifies memory sharing
 - ♦ *What is the essence of the mechanism that enables this?*
 - But, a user process had to be completely loaded into memory before it could run
- Wasteful since a process only needs a small amount of its total memory at any time (*reference locality!*)

Virtual Memory

- Definition: *Virtual memory* permits a process to run with only some of its virtual address space loaded into physical memory
- Key idea: Virtual address space translated to either
 - ◆ Physical memory (small, fast) or
 - ◆ Disk/SSD (backing store), large but slow
- Deep thinking – what made above possible?
- Objective:
 - ◆ To produce the illusion of memory as big as necessary

Virtual Memory

- “To produce the illusion of memory as big as necessary”
 - ◆ Without suffering a huge slowdown of execution
 - ◆ What makes this possible?
 - ◆ *Principle of locality*
 - » Knuth’s estimation of 90% of the time in 10% of the code
 - » There is also significant locality in data references

Virtual Memory Implementation


- Virtual memory is typically implemented via *demand paging*
- *demand paging*:
 - ♦ Load memory pages (from storage or initially allocated) “**on demand**”
 - ♦ paging with swapping, e.g., physical pages are swapped in and out of memory

Demand Paging

(paging with swapping)

- If not all of a program is loaded when running, what happens when referencing a byte not loaded yet?

- How to **detect** this?
 - ♦ In software?



Demand Paging

(paging with swapping)

- If not all of a program is loaded when running, what happens when referencing a byte not loaded yet?
- Hardware/software cooperate to make things work
 - ◆ Include a valid bit (present bit) in each PTE
 - ◆ Any page not in main memory right now has the valid bit cleared in its PTE
 - ◆ If valid bit isn't set, a reference to the page results in a trap by the paging hardware, called *page fault*
 - ◆ What needs to happen when page fault occurs?

What happens at virtual memory allocation time and access time?

- What happens at virtual memory allocation time?
 - ♦ If demand paging (on-demand allocation) is used, the OS allocates a virtual address (more later today) and establishes a PTE with no PFN and with invalid bit set
- What happens when the virtual address is first accessed?
 - ♦ The OS should allocate physical memory for it
 - ♦ How to capture the first write to a virtual page?
 - » e.g. want to trap into page fault handler
 - Use valid bit
 - ♦ In page fault handler handler, check if the virtual page is allocated (and access permitted)
 - » If not, segmentation fault
 - » Else allocate physical page and update PTE

What happens when main memory is not big enough?

- Processes running on a machine have collectively used more memory than what the physical main memory has.
- Some memory pages need to be put to storage (**swap out**)
- What happens when a swapped out page is accessed?
 - ♦ Need to **swap in** the page
 - ♦ How to detect that a swapped out page is accessed?

Next time...

- Chapter 22

Next time...

- Chapter 22